

How to (legally) keep secrets from mobile operators

Anonymized for submission

No Institute Given

Abstract. Secure-channel establishment allows two endpoints to communicate confidentially and authentically. Secure channels hide all the information sent across them, good or bad, to everyone but the endpoints; thus, they are often targeted by mass surveillance in the name of (inter)national security. While some protocols are constructed to allow easy mass interception of communications, others are designed to preserve user privacy and are either subverted or prohibited to use without trapdoors.

We introduce LIKE, a primitive that provides secure-channel establishment with an exceptional, session-specific opening mechanism. Designed for mobile communications, where an operator forwards messages between the endpoints, it can also be used in other settings. LIKE allows Alice and Bob to establish a secure channel with respect to n authorities. If the authorities all agree on the need for interception, they can ensure that the session key is retrieved. We show that LIKE is versatile with respect to who actually learns that key, and prove that the session key remains indistinguishable from random as long as at least one authority is honest and does not endorse interception. Furthermore, we guarantee non-frameability – nobody can incriminate a user of taking part in a conversation it hasn't actually had – and honest-operator: if the operator accepts a transcript as valid, then the key retrieved by the authorities is the key that should have been computed by Alice and Bob. Experimental results show that our protocol can be efficiently implemented.

1 Introduction

For almost a decade mass surveillance has caused controversy, widespread protests, and scandals; and yet, it is on the rise, affecting users every day. The NSA, for instance, illegally collected phone-call data from all Verizon customers, and the data of all calls occurring in the Bahamas and Afghanistan¹. During the present covid-19 pandemic, Germany used contact-tracing data to pursue criminal investigations². The need for privacy-enhancing solutions that provide transparency and limit mass surveillance has never been greater.

User privacy is a human right acknowledged, *e.g.*, by Article 12 of the Universal Declaration for Human Rights³: “No one shall be subjected to arbitrary interference with his privacy [...] or correspondence [...]. Everyone has the right to the protection of

¹ <https://mashable.com/2014/06/05/edward-snowden-revelations/?europa=true>

² <https://www.fairtrials.org/news/short-update-police-germany-defend-use-contact-tracing-criminal-investigations>

³ <https://www.un.org/en/universal-declaration-human-rights/>

the law against such [...] attacks.” The European General Data Protection Regulation (GDPR) and e-Privacy both aim to protect privacy in digital environments, requiring minimal, transparent, secure, and user-controlled storage of data. On the side of communication applications, our awareness of mass surveillance has pushed towards new encryption options, *e.g.*, in WhatsApp and viber. Data sent over mobile networks is also encrypted by mobile network operators, which have an incentive to offer users more privacy-protection than their competitors.

Unfortunately, mobile data remains at risk, especially when exceptional access to it lies with a single entity. This was the case when only the integrity of Tim Cook (Apple CEO) and his awareness of the danger of such a precedent prevented encrypted phone data to be given to the FBI⁴. Far from being a deterrent, privacy and encryption are under a bigger threat than ever⁵.

Law-enforcement agencies argue that mobile data can be pivotal to investigations, which are undermined by individuals ‘going dark’⁶. Even strong privacy advocates, such as Abelson *et al.* [6] agree that *targeted* investigations, *limited in scope and motivation*, can be legitimate and useful (see the corruption-scandal regarding Nicolas Sarkozy’s campaign funds). It is this type of limited LI that emerging EU legislation advocates⁷.

Recent research has tried to find a middle ground between privacy and lawful interception [37], by making lawful interception possible, but very costly. Unfortunately, this idea would not be adopted in practice, since it sacrifices the crucial timely (exceptional) decryption requirement [2]. In this paper we define and instantiate Lawful-Interception Key-Exchange (LIKE), a novel primitive allowing mobile users to E2E encrypt their conversation, but providing exceptional lawful interception. This would remove the “[...] need to choose between compliance and strong encryption” (in the words of Joel Wallenstrom).

LIKE combats mass surveillance in multiple ways. It *fine-grains* the window of interception to a single session, such that one exceptional opening will not affect past or future conversations. Moreover, the freshness used in the protocol is user- not authority-generated, thus *removing the need* for a centralized, secure party storing all the keys. We also divide the responsibility of exceptional opening between multiple authorities that *must agree* to lawfully intercept communications. Finally, we make the primitive *versatile*, allowing only one, or only several authorities to ultimately retrieve the session key. This makes our solution more privacy-preserving than mobile protocols today, while at the same time remaining compatible to the strictest LI-supporting 3GPP texts [2].

Lawful interception in standards. Regardless of one’s stand on it, Lawful Interception (LI) is a part of our world, regulated by laws and standards. Disregarding it can lead to

⁴ <https://www.wired.com/story/the-time-tim-cook-stood-his-ground-against-fbi/>

⁵ <https://www.europol.europa.eu/newsroom/news/europol-and-european-commission-inaugurate-new-decryption-platform-to-tackle-challenge-of-encrypted-material-for-law-enforcement>

⁶ <https://www.fbi.gov/news/speeches/going-dark-are-technology-privacy-and-public-safety-on-a-collision-course>

⁷ https://files.orf.at/vietnam2/files/fm4/202045/783284_fh_st12143-re01en20_783284.pdf

privacy threats (subversion and mass surveillance), and solutions that provide privacy but no LI are discarded as unlawful, regardless of their merits. We take the alternative approach: we analyze LI requirements and technically provide for them, while still maximizing data privacy.

LI requirements in mobile communications were put forth by 3GPP and adopted by standardization organizations such as ETSI. LI is defined to be the procedure through which a law enforcement agency, holding a legal warrant, can obtain information about phone calls – either metadata (time of calls, identity of the callers) or communication content (of conversations happening in real time). By law, a mobile network operator *must* provide the data requested and specified by a warrant. Three main types of requirements regulate LI: user-privacy requirements, LI-security requirements, and requirements on encryption.

User Privacy: the interception is limited in time (as dictated by the warrant) and to a targeted user. The law enforcement agency should not get data packages or conversations outside the warrant, from the same or other users.

LI Security: LI must be undetectable by either users (for whom the quality of service must remain the same) or non-authorized third parties (including other law-enforcement agencies). The intercepted data must be provided as soon as possible, with no undue delay.

Special case: if mobile network encryption is implemented, the operator must provide either the decrypted content, or a means to decrypt that content (*e.g.*, a decryption key). However, if the users employ other means of encryption, not provided by the operator, the latter is under no obligation to provide decrypted (or decryptable) conversations.

1.1 Our contributions

Based on these requirements, we define a novel cryptographic primitive called Lawful-Interception Key Exchange (LIKE, in short), for which we formalize the following strong security properties:

C Correctness: Under normal conditions, Alice and Bob obtain the same key. Moreover, this will be the key retrieved by the collaboration of all the authorities by means of lawful interception (requirement R2)

KS Key-security: If at least one authority and both users are honest for a given session, that session's key remains indistinguishable from a random key of the same length with respect to an adversary that can control all the remaining parties (including the other authorities and the operator); (requirements R1 and R2)

NF Non-Frameability: The collusion of malicious users, the authorities, and the operator cannot frame an honest user of participating to a session she has not been a party to (requirement R2);

HO Honest Operator: If an honest operator forwards the so-called session state (see below) of a session it deems correct, then the key recovered by the authorities is the one that the session transcript should have yielded (requirement R3). Thus operators can prove that this protocol is compliant with LI specifications.

In LIKE, after setup and key-generation, mobile users like Alice and Bob can run an authenticated key-exchange (AKE) protocol in the presence of their respective oper-

ators. This protocol allows (only) the end users to compute session keys; the operators output a public value called a session state.

We consider a set of n authorities entitled to performing LI within the limits of the law (authorities can be, *e.g.*, a court of justice, a law enforcement agency, or operators). Given the session state, authorities may each extract a token from that state, and the use of *all* the n tokens can yield the key. Importantly, unless they are an authority, operators *cannot* recover Alice and Bob’s key; instead they forward and verify the compliance of exchanged messages (else, the operator halts communication, believing Alice and Bob are misbehaving).

Our protocol. As our second contribution, we describe an instantiation of LIKE using standard building blocks (signatures, zero knowledge proofs and signatures of knowledge) and we prove the security of our scheme based on the hardness of the Bilinear Decisional Diffie-Hellman problem, the unforgeability of our signature scheme and the security of our proofs and signatures of knowledge.

Mindful of the practical requirements of the protocol, we place most of the storage and computational burden during AKE on the operator (not on the endpoints). Our proofs and signatures of knowledge are simple and can be implemented based on Schnorr and respectively Chaum and Pedersen proofs (with Fiat-Shamir). The two endpoints do have to compute a pairing operation – however, we explain that the actual computation can actually be delegated to the mobile phone, leaving only a single exponentiation (in the target group) to be performed on the USIM card.

The complexity of the opening procedure is reduced, but linear in the number of authorities. Some steps, such as the token generation, are parallelizable, but others (such as putting those tokens together) are not. However, even so the computational burden remains minimal and would not contradict requirements R2 (since all computations are done by the authorities, not the operator, the quality of service is the same) and R3 (no undue delay would be incurred by the operator, and only minimal delays occur at the authorities). A proof-of-concept implementation given in Section 7 illustrates this point.

1.2 Related work

Existing encryption (through the AKA protocol) in mobile networks is not E2E secure, only providing privacy with respect to non-authorized third parties. AKA is compliant to LI requirements [3–5] because the secure channel it provides is between the user and the operator (rather than user-to-user); thus the operator has unrestricted access to all user communications. Our LIKE protocol provides much stronger privacy in that respect.

LIKE also provides much stronger guarantees than key-escrow [30, 19, 7, 36, 33, 9, 17, 27, 25, 26, 21, 28, 34, 31, 18]: we require no trust in authorities with respect to their input; we fine-grain exceptional opening so that it only holds for one session at a time; we allow authorities to remain offline at all times except for exceptional opening; we minimize storage and computational costs for users and authorities; we have no central key-generation authority which knows all secret keys; and we guarantee the new properties of non-frameability and honest operator – which are tailored to the LI requirements analyzed above.

Our work is motivated by the same problem that has yielded works such as [10, 37], but we take a different approach: instead of making LI computationally costly, we instead limit its scope, divide the power to act between several authorities, and fine-grain access. Our setup also resembles that of reverse firewalls [29] – which builds on related work pioneered by Young and Yung [38, 20, 23]. Although apparently similar, the setup of these works is complementary to ours. Kleptography describes ways for users to abuse protocols such that the latter appear to be running normally, while in fact the permits subliminal information to leak. For instance in key-exchange, a government agency could want to substitute the implementation of the protocol by one that has a backdoor, to make the key recoverable even without formally doing LI. We do not address this problem here; instead, we treat the case in which Alice and Bob pick their input so as to make it hard for the authorities to retrieve the key *even when LI is correctly triggered* (the HO property).

2 Preliminaries

The notation $x \leftarrow y$ indicates that the variable x takes a value y and $x \xleftarrow{\$} X$ indicates that the variable x is chosen from the uniform distribution on X . We write $A(x) \rightarrow a$ to express that the algorithm A , running on input x , outputs a , and $P\langle A(x), B(y) \rangle(z) \rightarrow (a, b)$ to express that the protocol P implements the interactions of $A(x) \rightarrow a$ and $B(y) \rightarrow b$, where z is an additional public input of A and B . Let λ be a security parameter.

Definition 1 (BDDH assumption [13]). Let $\mathbb{G}_1 = \langle g_1 \rangle$, $\mathbb{G}_2 = \langle g_2 \rangle$, and \mathbb{G}_T be groups of prime order p of length λ . Let $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T$ be a type 3 bilinear map. The Bilinear decisional Diffie-Hellman problem (BDDH) assumption holds in $(\mathbb{G}_1, \mathbb{G}_2, \mathbb{G}_T, e)$ if, given $(a, b, c, d_1) \xleftarrow{\$} (\mathbb{Z}_p^*)^4$, $d_0 \leftarrow abc$, and $\beta \xleftarrow{\$} \{0, 1\}$, no PPT adversary \mathcal{A} can guess b from $(g_1^a, g_2^a, g_1^b, g_2^b, g_1^c, g_2^c, e(g_1, g_2)^{d_\beta})$ with non-negligible advantage. We denote by $\text{Adv}^{\text{BDDH}}(\lambda)$ the maximum advantage over all PPT adversaries.

Definition 2 (Digital signatures). A digital signature scheme $\text{DS} = (\text{SGen}, \text{SSig}, \text{SVer})$ is defined by three algorithms: $\text{SGen}(1^\lambda) \rightarrow (\text{PK}, \text{SK})$; $\text{SSig}(\text{SK}, m) \rightarrow \sigma$, and $\text{SVer}(\text{PK}, m, \sigma) \rightarrow b$. Let Sig be an oracle that, on input m , returns $\text{SSig}(\text{SK}, m)$. Let \mathcal{A} be PPT adversary.

We define the following experiment:

$\text{Exp}_{\text{DS}}^{\text{EUF-CMA}}(\mathcal{A})$:
 $(\text{PK}, \text{SK}) \leftarrow \text{SGen}(1^\lambda)$; $(m, \sigma) \leftarrow \mathcal{A}^{\text{Sig}(\cdot)}(\text{PK})$;
 Return 1 if (m, σ) was not output by $\text{Sig}(\cdot)$ and
 $\text{SVer}(\text{PK}, m, \sigma) = 1$, 0 otherwise.

A digital signature scheme DS is existentially unforgeable against chosen message attacks (EUF-CMA) if, for all PPT adversaries \mathcal{A} , $\text{Adv}_{\text{DS}, \mathcal{A}}^{\text{EUF-CMA}}(\lambda) = \mathbb{P} \left[\text{Exp}_{\text{DS}, \mathcal{A}}^{\text{EUF-CMA}}(\lambda) = 1 \right]$ is negligible in λ . We denote by $\text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda)$ the maximum advantage over all PPT adversaries.

Signature of Knowledge. Let \mathcal{R} be a binary relation and let \mathcal{L} be a language such that $s \in \mathcal{L} \Leftrightarrow (\exists w, (s, w) \in \mathcal{R})$. A Non-Interactive Proof of Knowledge (NIPoK) [8] allows

a prover to convince a verifier that he knows a witness w such that $(s, w) \in \mathcal{R}$. Here, we follow [14] and write $\text{NIPoK} \{w : (w, s) \in \mathcal{R}\}$ for the proof of knowledge of w for the statement s and the relation \mathcal{R} . A *signature of knowledge* essentially allows one to sign a message and prove in zero-knowledge that a particular statement holds for the key [15]. In this paradigm, w is a secret key and s is the corresponding public key.

Definition 3 (Signature of Knowledge). Let \mathcal{R} be a binary relation and \mathcal{L} be a language such that $s \in \mathcal{L} \Leftrightarrow (\exists w, (s, w) \in \mathcal{R})$. A Signature of Knowledge for \mathcal{L} is a pair of algorithms $(\text{SoK}, \text{SoKver})$ with $\text{SoK}_m \{w : (s, w) \in \mathcal{R}\} \rightarrow \pi$ and $\text{SoKver}(m, s, \pi) \rightarrow b$, such that:

- *Perfect Zero Knowledge:* There exists a polynomial time algorithm Sim , the simulator, such that $\text{Sim}(m, s)$ and $\text{SoK}_m \{w : (s, w) \in \mathcal{R}\}$ follow the same probability distribution.
- *Knowledge Extractor:* There exists a PPT knowledge extractor Ext and a negligible function ϵ_{SoK} such that for any algorithm $\mathcal{A}^{\text{Sim}(\cdot, \cdot)}(\lambda)$ having access to a simulator that forges signatures for chosen instance/message tuples and that outputs a fresh tuple (s, π, m) with $\text{SoKver}(m, s, \pi) = 1$, the extractor $\text{Ext}^{\mathcal{A}}(\lambda)$ outputs w such that $(s, w) \in \mathcal{R}$ having access to $\mathcal{A}(\lambda)$ with probability at least $1 - \epsilon_{\text{SoK}}(\lambda)$.

We omit to recall the definition of NIPoK which is the same as SoK without the messages.

3 LIKE protocols

We define lawful-interception (authenticated) key-exchange (LIKE) in terms of two mechanisms: a three-party AKE protocol featuring two mobile subscribers (which we call users) and a mobile network operator, and an Extract-and-Open mechanism involving a number of parties (which we call authorities).

Intuition. Our AKE component allows user Alice, subscribing to operator O_A , and Bob, subscribing to O_B , to compute a session key in the presence of those operators. However, O_A and O_B will not be able to compute the session key. Instead, they output some auxiliary material which we call *session state*, the latter allowing for session authentication and an ulterior key-recovery by a set of authorities.

In the Extract-and-Open component, each authority uses its secret key to *extract* a trapdoor from the operator's session state. Subsequently the trapdoors are used together to *open* the session key.

Formalization. Let USERS be a set of mobile users and OPS be a set of operators, such that each user is affiliated to precisely one operator. In particular, we equate users with uniquely affiliated hardware, such as SIM cards, rather than with people. We also consider a set of authorities AUTH of cardinality n , with elements indexed as A_1, \dots, A_n .

Let PARTIES be the set of all participants: $\text{USERS} \cup \text{OPS} \cup \text{AUTH}$. Mobile users have no super-role (they can be neither authorities, nor operators, $\text{USERS} \cap \text{AUTH} = \emptyset = \text{USERS} \cap \text{OPS}$). The case of operators being authorities is more complicated, and we assume that $\text{OPS} \cap \text{AUTH} = \emptyset$.

Although we formally introduce parties, attributes, and oracles in the next section, we anticipate a few notations here. For each party P we use a dot notation to refer to *attributes* of that party (such as a long-term private or public key). For instance $A.PK$ refers to Alice's public key and $\Lambda_i.SK$ refers to the secret key of the i^{th} authority. We slightly abuse notation and write O_A to indicate Alice's operator – in reality, in our scheme we abstract the process of registering to an operator by allowing two parties to run the protocol in the presence of any two operators they choose to have. We also assume that Alice and Bob always interact with their respective operators during the protocol, which can be easily achieved by using, *e.g.*, the standard AKA protocol as an overlay of ours.

Definition 4. A lawful interception key exchange protocol (LIKE) is defined by the following algorithms:

- $\text{Setup}(1^\lambda) \rightarrow pp$: Takes as input a security parameter in unary notation, and outputs pp , the public parameters of the system. They become known to all parties.
- $\text{UKeyGen}(pp) \rightarrow (U.PK, U.SK)$: Takes as input the public parameters pp and outputs a user key pair.
- $\text{OKeyGen}(pp) \rightarrow (O.PK, O.SK)$: Takes as input the public parameters pp and outputs the operator key pair.
- $\text{AKeyGen}(pp) \rightarrow (\Lambda.PK, \Lambda.SK)$: Takes as input the public parameters pp and outputs an authority key pair.
- $\text{AKE}(A(A.SK), O_A(O_A.SK), O_B(O_B.SK), B(B.SK)) \langle PK_{A \rightarrow B} \rangle \rightarrow (k_A, sst_A, sst_B, k_B)$: An authenticated key-exchange protocol between two subscribers $(A, B) \in \text{USERS}^2$, their operators $(O_A, O_B) \in \text{OPS}^2$, such that O_A and O_B provide active middleware for A and B during the protocol at all times (A and B never interact directly). Alice, Bob, and their operators each take as input their own secret key. In addition, each party has access to the set of public parameters $PK_{A \rightarrow B}$ which contains: the public parameters pp , the public keys of Alice and Bob ($A.PK, B.PK$), and a vector of authority public keys $APK = (\Lambda_i.PK)_{i=1}^n$ with $\Lambda_i \in \text{AUTH}$ for all i . At the end of the protocol, A (resp. B) returns a session secret key k_A (resp. k_B) and the operator O_A (resp. O_B) returns a (public) session state sst_A (resp. sst_B). In case of failure, the parties output a special symbol \perp instead.
- $\text{Verify}(pp, sst, A.PK, B.PK, O.PK, APK) \rightarrow b$: Takes as input a session state sst , two user public keys $A.PK$ and $B.PK$, the public key of an operator $O.PK$, a set of public keys of the authorities $APK = (\Lambda_i.PK)_{i=1}^n$, outputting a bit $b = 1$ if the sst was correctly generated and authenticated by O , and $b = 0$ otherwise.
- $\text{TDGen}(pp, \Lambda.SK, sst) \rightarrow \Lambda.t$: Takes as input an authority secret key $\Lambda.SK$ and session state sst , and outputs a trapdoor $\Lambda.t$.
- $\text{Open}(pp, sst, APK, \mathcal{T}) \rightarrow k$: Takes as input session state sst , a vector of public keys of the authorities $APK = (\Lambda_i.PK)_{i=1}^n$, a vector of the corresponding trapdoors $\mathcal{T} = (\Lambda_i.t)_{i=1}^n$, and outputs a session secret key k . In case of failure, k may take a special value \perp .

Definition 5 (Correctness). Let λ a security parameter and n an integer. Run $pp \leftarrow \text{Setup}(1^\lambda)$, $(A.PK, A.SK) \leftarrow \text{UKeyGen}(pp)$, $(B.PK, B.SK) \leftarrow \text{UKeyGen}(pp)$, $(O_A.PK, O_A.SK) \leftarrow \text{OKeyGen}(pp)$, $(O_B.PK, O_B.SK) \leftarrow \text{OKeyGen}(pp)$. For all $i \in \llbracket 1, n \rrbracket$, $(\Lambda_i.PK, \Lambda_i.SK) \leftarrow \text{AKeyGen}(pp)$. Let $APK \leftarrow (\Lambda_i.PK)_{i=1}^n$. Then:

- $PK_{A \rightarrow B} \leftarrow (pp, A.PK, B.PK, APK);$
- $(k_A, sst_A, sst_B, k_B) \leftarrow$
 $AKE\langle A(A.SK), O_A(O_A.SK), O_B(O_B.SK), B(B.SK) \rangle(PK_{A \rightarrow B});$
- $b_A \leftarrow \text{Verify}(pp, sst_A, A.PK, B.PK, O_A.PK, APK);$
- *For all i in $\llbracket 1, n \rrbracket$, $\Lambda_i.t_A \leftarrow \text{TDGen}(pp, \Lambda_i.SK, sst_A);$*
- $k_A^* \leftarrow \text{Open}(pp, sst_A, (\Lambda_i.PK)_{i=1}^n, (\Lambda_i.t_A)_{i=1}^n);$
- $b_B \leftarrow \text{Verify}(pp, sst_B, A.PK, B.PK, O_B.PK, APK);$
- *For all i in $\llbracket 1, n \rrbracket$, $\Lambda_i.t_B \leftarrow \text{TDGen}(pp, \Lambda_i.SK, sst_B);$*
- $k_B^* \leftarrow \text{Open}(pp, sst_B, APK, (\Lambda_i.t_B)_{i=1}^n).$

For any $(b_A, b_B, k_A, k_A^, k_B, k_B^*)$ generated as above: $\Pr[b_A = b_B = 1 \wedge k_A = k_A^* = k_B = k_B^*] = 1.$*

To use a lawful-interception key-exchange scheme, users like Alice and Bob generate their keys using UKeyGen, operators generate their keys using OKeyGen, and each authority generates its keys using AKeyGen. Then, Alice, Bob, and the operators run the protocol AKE using the vector of all the authority public keys. At the end of the protocol, Alice and Bob share a session secret key (not known by the operator) that they can use to communicate securely. Each operator returns a public session state sst , which is heavily protocol-dependent: its main purpose is to encapsulate authenticated session data that will eventually allow the authorities to verify that the protocol was run correctly, and recover the session key. Finally, any party can check, by using the verification algorithm Verify , that the session secret key was honestly generated from a valid execution of the protocol between Alice, Bob and the operators.

The authorities may later retrieve this session secret key. Each authority uses its secret key and the session state to compute a trapdoor to that key. Before computing it, the authority checks the soundness of sst using Verify . By using all the trapdoors and the algorithm Open , one retrieves Alice and Bob's session key.

Depending on the lawful interception scenario, the Open algorithm may be run by one or multiple parties (we only require that whoever runs the algorithm possesses all the trapdoors). Our protocol is versatile and can adapt to many cases.

4 Security Model

We proceed to detail the adversarial model and formal security properties required for our new primitive LIKE. This is an essential contribution of our paper, for the following two reasons: first, because a formal treatment of such schemes allows us to *prove and quantify* the security of our scheme; and secondly, because the properties we formalize are much stronger than the obvious extension of the key-security properties required in regular authenticated key exchange. In particular, the properties of honest operator and non-frameability, for instance, are not achievable by schemes relying on a central key-distribution authority, like [34, 31, 18].

The rest of this section is structured as follows. First we give the adversarial model, describing how parties and protocol instances are run. Then we list the oracles which adversaries can use to manipulate the honest parties. Finally, we describe formal security games for each of the given properties.

The adversarial model Each of the parties mentioned in Section 3 (the users, operators, and authorities) has a unique role, which is assumed not to change during the course of our security experiments. Each party P is associated with a number of attributes, listed below.

- (SK, PK): a tuple consisting of a long-term private key SK and a public key PK. They may be instantiated to values output by UKeyGen, OKeyGen, or AKeyGen.
- γ : a corruption flag with a bit value, which indicates whether that party has been corrupted by the adversary or not. This bit is initially set to 0, but may be turned to 1 during the security game. Afterwards, the bit may no longer be changed to 0.

We continue to use the dot notation in Section 3; thus, *e.g.*, an operator’s corrupt bit is denoted by $O.\gamma$.

Alice, Bob, and their operators run the AKE in sessions. At each new protocol execution, a new *instance* of that party is created. Four party instances (one each for Alice, her operator, Bob’s operator, and Bob himself) run the protocol together to form a session. We denote by π_P^i the i -th instance of party P .

Instances of each party automatically have access to the values of the long-term keys of those parties and their corruption bit. In addition, they keep track of the following attributes.

- sid: a session identifier consisting of a tuple of values (session specific), such as portions of the public parameters and the randomness. We stress that this attribute stores *only* the state information pertinent to that protocol. This value is internally computable by the parties interacting in the protocol, but without involving any secret value. This attribute is instantiated to a default value \perp and changes its value during the protocol’s run.
- PID: partner identifiers. If $P \in \text{USERS}$, then $\text{PID} \in \text{USERS}$ such that $P \neq \text{PID}$, else $P \in \text{OPS}$ and $\text{PID} \in \text{USERS}^2$ such that the two elements of the vector PID are different.
- OID: operator identifiers. If $P \in \text{USERS}$ then $\text{OID} \in \text{OPS}^2$ such that $|\text{OID}| = 2$ (the operators are distinct). Else $\text{OID} \in \text{OPS}$.
- AID: authority identifiers such that $\text{AID} \in \text{AUTH}^n$.
- α : the instance’s accept flag, which is undefined ($\alpha = \perp$) until the instance terminates its AKE execution. If the instance terminates by aborting the protocol, it sets $\alpha = 0$, else (protocol terminates without error), $\alpha = 1$.
- k : the instance’s session key, first initialized to a special value \perp . If the protocol terminates without error, k takes the value returned by the protocol. An operator instance does not have this attribute (it is replaced by the session state attribute sst defined below).
- sst: a set of variables storing additional state of the instance’s protocol execution. If the protocol terminates without error, sst takes the value returned by the protocol. User instances do not have this attribute.
- ρ : the instance’s reveal bit, first initialized to 0 and set to 1 if the adversary reveals a session key $k \neq \perp$. Operator instances do not have this attribute.
- b : a bit chosen uniformly at random upon the creation of the instance.
- τ : the transcript of the ongoing session, first initialized as \perp , turning to the ordered list of messages sent and received by that instance in the same order.

We define an auxiliary function $\text{IdentifySession}(\text{sst}, \pi)$ that takes as input a session state sst and a party instance π , and outputs 1 if π took part in the session where sst was created, and 0 otherwise. This is because during the session opening, the authorities must be able to extract and verify the session identifier for themselves.

Notice that, unlike in typical authenticated key-exchange protocols, we have two different kind of parties in this protocol (users and operators), and three types of parties that partner an instance (as indicated by the attributes PID, OID, and AID). The instance stores mobile user partners in PID, operator partners in OID, and authorities partnering it in AID.

Moreover, IdentifySession and sst are protocol dependent, *i.e.* they will have a different instantiation depending on the protocol we analyse. For our AKE component we require the notion of matching conversation defined as follows:

Definition 6 (Matching instances). *For any $(i, j) \in \mathbb{N}^2$ and $(A, B) \in \text{USERS}^2$ such that $A \neq B$, we say that π_A^i and π_B^j have matching conversation if all the following conditions hold: $\pi_A^i.\text{sid} \neq \perp$, $\pi_A^i.\text{sid} = \pi_B^j.\text{sid}$, and $\pi_A^i.\text{AID} = \pi_B^j.\text{AID}$. If two instances π_A^i and π_B^j have matching conversation, we sometimes say, by abuse of language, that π_A^i matches π_B^j .*

Oracles. We define the security properties of our novel LIKE primitive in terms of games that an adversary plays against a challenger (simulating all the honest parties). In order to manipulate the environment, the adversary will have access to some or all of the oracles presented below. Intuitively, we give the adversary the ability to register honest or malicious participants (the Register oracle), initiate new sessions (the NewSession oracle), interact in the AKE protocol (the Send oracle), corrupt parties (the Corrupt oracle), reveal session keys (the Reveal oracle), or reveal trapdoors towards potentially opening the key by the LI process (the RevealTD oracle). Finally, we use a testing oracle (Test), which will test whether the session key remains indistinguishable from random from the adversary's point of view.

For each oracle, we add that they abort if they receive a query that is not well formatted, or if the challenger does not have enough information to correctly answer the query.

- $\text{Register}(P, \text{role}, \text{PK}) \rightarrow \perp \cup P.\text{PK}$: On input a party identity $P \notin \text{USERS} \cup \text{OPS} \cup \text{AUTH}$, a role $\text{role} \in \{\text{user}, \text{operator}, \text{authority}\}$ and a public key PK :
 - If $\text{role} = \text{user}$ (resp. operator, authority), the oracle adds P to the set USERS (resp. OPS , AUTH).
 - If $\text{role} = \text{user}$ (resp. operator and authority) and $\text{PK} = \perp$, then it runs $\text{UKeyGen}(\text{pp}) \rightarrow (P.\text{PK}, P.\text{SK})$ (resp. OKeyGen and AKeyGen).
 - If $\text{PK} \neq \perp$, it sets the corruption bit γ to 1 for this party, sets $P.\text{PK} = \text{PK}$ and $P.\text{SK} = \perp$.

Finally, it returns $P.\text{PK}$.

- $\text{NewSession}(P, \text{PID}, \text{OID}, \text{AID}) \rightarrow \pi_P^i$: On input a party $P \in \text{USERS} \cup \text{OPS}$, if $P \in \text{USERS}$ then PID , OID and AID must verify that $\text{PID} \in \text{USERS}$ such that $P \neq \text{PID}$, $\text{OID} \in \text{OPS}^2$ and $\text{AID} \in \text{AUTH}^*$ such that $|\text{AID}| \neq 0$. If $P \in \text{OPS}$ then PID , OID and AID verify that $\text{PID} \in \text{USERS}^2$ such that PID contains two different

- users, $\text{OID} \in \text{OPS}$ and $\text{AID} \in \text{AUTH}^*$ such that $|\text{AID}| \neq 0$. On the i^{th} call to this oracle, it returns a new instance π_P^i with the attributes PID, OID and AID.
- $\text{Send}(\pi_P^i, m) \rightarrow m'$: Sends the message m to the instance π_P^i and returns a new message m' , as specified by the protocol AKE. This message can be a special value \perp in case of failure (if the instance does not exist, has already rejected the session, the session is finished, the protocol aborts, m is not well formatted, or $\text{P.SK} = \perp$).
 - $\text{Reveal}(\pi_P^i) \rightarrow k$: For an instance of $P \in \text{USERS}$ that has accepted the session ($\alpha = 1$), it returns the session key $\pi_P^i.k$ and sets the reveal bit $\pi_P^i.\rho$ to 1. For an instance of $P \in \text{OPS}$ that has accepted the session ($\alpha = 1$), it returns the session state $\pi_P^i.\text{sst}$ and sets the reveal bit $\pi_P^i.\rho$ to 1. If the session has not been accepted ($\alpha \neq 1$), the oracle returns \perp .
 - $\text{Corrupt}(P) \rightarrow \text{P.SK}$: It returns the SK of the party $P \in \text{USERS} \cup \text{OPS} \cup \text{AUTH}$ and sets the corruption bit $\text{P}.\gamma$ to 1 for this party and any of its instances.
 - $\text{Test}(\pi_P^i) \rightarrow \tilde{k}$: Can be queried for an instance of $P \in \text{USERS}$ that has accepted the session ($\alpha = 1$). If $\pi_P^i.b = 0$, it returns the session key $\pi_P^i.k$. Otherwise, it returns a randomly sampled value r from the same domain as $\pi_P^i.k$. If the instance has not accepted the session, it returns \perp . If this oracle had been previously queried and had returned a value different from \perp , it will return \perp (this oracle can effectively be queried only once).
 - $\text{RevealTD}(\text{sst}, A, B, O, (\Lambda_i)_{i=1}^n, l) \rightarrow \Lambda_l.t$: If $\text{Verify}(\text{pp}, \text{sst}, A.\text{PK}, B.\text{PK}, O.\text{PK}, (\Lambda_i.\text{PK})_{i=1}^n) = 1$, then it runs $\Lambda_l.t \leftarrow \text{TDGen}(\text{pp}, \Lambda_l.\text{SK}, \text{sst})$ and returns $\Lambda_l.t$, else it returns \perp .

We emphasize that the operators involved in the session do not contribute to the security of the key in the traditional AKE sense. Instead, operators verify the soundness of the exchanges and produce a session state sst ; this value then serves to prove to the authorities that the operator outputting correctly arbitrated the protocol from its point of view. We do not require the operators to agree on sst , and perform the opening procedure on a single operator at a time (since this is what would happen in most real-life situations). If one operator validates, but the other operator aborts the protocol, then ultimately the session will not take place.

Security games. We define the security for LIKE protocols in terms of three properties: key-security (KS), non-frameability (NF), and honest operator (HO).

Key-security. Our KS game is an extension of the standard notion of AKE security [11]. The adversary has access to all the oracles presented above (subject to the definition of key-freshness defined below). Its goal is to distinguish the real key used by Alice and Bob from a key picked randomly from the same domain.

We give the adversary considerably more power than in traditional 2-party AKE games. The attacker can adaptively corrupt all but the two users targeted in the challenge, all the operators, and all but one of the authorities. The adversary may also register illegitimate users and learn all but one of the trapdoors involved in any session.

In other words, this definition only allows the key to be known by the parties having computed it (Alice and Bob), and by the collusion of *all* the authorities (LI). We rule out these two circumstances by the following notion of key freshness. All other (collusions of) parties should fail to distinguish the real session key from a random one.

$\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda)$:
 $\text{pp} \leftarrow \text{Setup}(1^\lambda)$;
 $\mathcal{O}_{\text{KS}} \leftarrow \left\{ \begin{array}{l} \text{Register}(\cdot, \cdot, \cdot), \text{NewSession}(\cdot, \cdot, \cdot), \text{Send}(\cdot, \cdot), \\ \text{Reveal}(\cdot), \text{RevealTD}(\cdot, \cdot), \text{Corrupt}(\cdot, \cdot), \text{Test}(\cdot) \end{array} \right\}$;
 $(i, \text{P}, \text{d}) \leftarrow \mathcal{A}^{\mathcal{O}_{\text{KS}}}(\lambda, \text{pp})$;
 If $\pi_{\text{P}}^i.k$ is fresh and $\pi_{\text{P}}^i.b = \text{d}$, then return 1;
 Else $b' \xleftarrow{\$} \{0, 1\}$, return b' .

Table 1: The key-security (KS) experiment.

Definition 7 (Key freshness). Let π_{P}^j be the j -th instance of a party $\text{P} \in \text{USERS}$ and denote by \mathcal{A} a PPT adversary against a LIKE scheme. We parse $\pi_{\text{P}}^j.\text{PID}$ as P' and $\pi_{\text{P}}^j.\text{AID}$ as $(\Lambda_i)_{i=1}^n$. We say that $\pi_{\text{P}}^j.k$ is fresh if all the following conditions hold:

- P has accepted in session j ($\pi_{\text{P}}^j.\alpha = 1$), P remained uncorrupted ($\text{P}.\gamma = 0$) up to the moment when $\pi_{\text{P}}^j.\alpha$ became 1, and the session key remains unrevealed ($\pi_{\text{P}}^j.\rho = 0$).
- if π_{P}^j matches some $\pi_{\text{P}'}^k$, for $k \in \mathbb{N}$, then P' has accepted in session k ($\pi_{\text{P}'}^k.\alpha = 1$), P' was uncorrupted ($\text{P}'.\gamma = 0$) up to when $\pi_{\text{P}'}^k.\alpha$ became 1, and the session key remains unrevealed ($\pi_{\text{P}'}^k.\rho = 0$).
- if no $\pi_{\text{P}'}^k$ matches π_{P}^j , P' is uncorrupted ($\text{P}'.\gamma = 0$).
- there exists $l \in \llbracket 1, n \rrbracket$ such that \mathcal{A} has never queried RevealTD on input $(\text{sst}, \text{A}, \text{B}, (\Lambda'_i)_{i=1}^{n'}, l')$ such that:
 - Λ_l is uncorrupted ($\Lambda_l.\gamma = 0$) and $\Lambda_l = \Lambda'_l$;
 - $\text{IdentifySession}(\text{sst}, \pi_{\text{P}}^j) = 1$.

In this game, the adversary eventually outputs an instance for which it guesses its test bit b . This instance must be key-fresh with respect to Definition 7. More formally, we consider the following key-security game:

Definition 8. We define the advantage of an adversary \mathcal{A} in the $\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda)$ experiment in Table 1 as:

$$\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda) := \left| \mathbb{P} \left[\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda) = 1 \right] - \frac{1}{2} \right|.$$

A lawful interception authenticated key-exchange scheme LIKE is key-secure if for all PPT adversaries \mathcal{A} , $\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda)$ is negligible as a function of the security parameter λ .

Non-Frameability. In this game, the adversary attempts to frame a target user P^* of taking part in an AKE session (corresponding to session state sst) in which P^* never took part or which P^* never accepted. The adversary is allowed to corrupt all parties in the system apart from P^* . More formally, we consider the following security experiment.

Definition 9. The advantage of an adversary \mathcal{A} in the non-frameability experiment $\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda)$ in Table 2 is defined as:

$$\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda) = \mathbb{P} \left[\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda) = 1 \right].$$

$\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda)$:
 $\text{pp} \leftarrow \text{Setup}(1^\lambda)$;
 $\mathcal{O}_{\text{NF}} \leftarrow \left\{ \begin{array}{l} \text{Register}(\cdot, \cdot, \cdot), \text{NewSession}(\cdot, \cdot, \cdot), \text{Send}(\cdot, \cdot), \\ \text{Reveal}(\cdot), \text{RevealTD}(\cdot, \cdot), \text{Corrupt}(\cdot, \cdot) \end{array} \right\}$;
 $(\text{sst}, \text{P}) \leftarrow \mathcal{A}^{\mathcal{O}_{\text{NF}}}(\lambda, \text{pp})$;
 If $\exists (A, B) \in \text{USERS}^2, n \in \mathbb{N}, O \in \text{OPS}$ and $(A_i)_{i=1}^n \in \text{AUTH}^n$ s.t.:
 $\text{Verify}(\text{pp}, \text{sst}, A.\text{PK}, B.\text{PK}, O.\text{PK}, (A_i.\text{PK})_{i=1}^n) = 1$ and
 $P \in \{A, B\}$ and
 $P.\gamma = 0$ and
 $\forall i, \text{if } \pi_P^i \neq \perp \text{ then } \text{IdentifySession}(\text{sst}, \pi_P^i) = 0 \text{ or } \pi_P^i.\alpha = 0$,
 Then return 1,
 Else return 0.

Table 2: The non-frameability (NF) experiment.

$\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda)$:
 $\text{pp} \leftarrow \text{Setup}(1^\lambda)$;
 $\mathcal{O}_{\text{HO}} \leftarrow \left\{ \begin{array}{l} \text{Register}(\cdot, \cdot, \cdot), \text{NewSession}(\cdot, \cdot, \cdot), \text{Send}(\cdot, \cdot), \\ \text{Reveal}(\cdot), \text{RevealTD}(\cdot, \cdot), \text{Corrupt}(\cdot, \cdot) \end{array} \right\}$;
 $(j, \text{sst}, A, B, O, (A_i, A_i.t)_{i=1}^n) \leftarrow \mathcal{A}^{\mathcal{O}_{\text{HO}}}(\lambda, \text{pp})$;
 If $O.\gamma = 1$ then return \perp ;
 If $\text{Verify}(\text{pp}, \text{sst}, A.\text{PK}, B.\text{PK}, O.\text{PK}, (A_i.\text{PK})_{i=1}^n) = 0$ then return \perp ;
 If $\text{IdentifySession}(\text{sst}, \pi_O^j.\text{sid}) = 0$ then return \perp ;
 $k_* \leftarrow \text{Open}(\text{pp}, \text{sst}, (A_i.\text{PK})_{i=1}^n, (A_i.t)_{i=1}^n)$;
 Return $(k_*, \pi_O^j, \{P_i.\text{PK}\}_{i=1}^{q_r})$.

Table 3: The honest-operator HO experiment. Here, q_r denotes the number of queries to the Register oracle, and P_i denotes the party input as the i -th such query.

A lawful interception authenticated key-exchange scheme LIKE is non-frameable if for all PPT adversaries \mathcal{A} , $\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda)$ is negligible as a function of the security parameter λ .

Honest Operator. The HO game captures the fact that honest operators have the power of aborting sessions if the messages exchanged are not well-formed or authentic. The adversary aims to forge a valid session state for which the operators have not aborted, but for which the authorities would Open to another key than that computed by Alice and Bob. The attacker can corrupt all the parties except for the operator which is assumed to eventually approve the session state, and it can even provide trapdoors.

Ideally, we would like LIKE schemes to guarantee that if the (honest) operator deems a session to be correctly run, then the extracted key (output by the experiment in Table 3) will be the one used by Alice and Bob. However, malicious participants Alice and Bob could run a protocol perfectly and then use a different key (which they might exchange out of band) to encrypt messages. This is inherent to secure-channel establishment when the participants are malicious. The best an LIKE protocol can guarantee is that the extracted key is the one that *would have resulted* from an honest run of the protocol whose session state is given. We express this in terms of a *key extractor*.

Intuitively, given as input an operator instance and a set of public keys, the extractor outputs the key k that should have been output by the partnered instances running that session, if they were honest.

Definition 10 (Key extractor). For any LIKE, a key extractor $\text{Extract}(\cdot, \cdot)$ is a deterministic unbounded algorithm such that, for any users A and B , operators O_A and O_B , and set of n authorities $(\Lambda_i)_{i=1}^n$, any set $\{\text{pp}, A.PK, A.SK, B.PK, B.SK, O_A.PK, O_A.SK, O_B.PK, O_B.SK, k, \text{sst}, \text{APK} = (\Lambda_i.PK)_{i=1}^n, (\Lambda_i.SK)_{i=1}^n, \tau_A, \tau_B\}$ generated as follows:

$\text{pp} \leftarrow \text{Setup}(\lambda); (A.PK, A.SK) \leftarrow \text{UKeyGen}(\text{pp});$
 $(B.PK, B.SK) \leftarrow \text{UKeyGen}(\text{pp});$
 $(O_A.PK, O_A.SK) \leftarrow \text{OKeyGen}(\text{pp});$
 $(O_B.PK, O_B.SK) \leftarrow \text{OKeyGen}(\text{pp});$
 For all $i \in [1, n]$, $(\Lambda_i.PK, \Lambda_i.SK) \leftarrow \text{AKeyGen}(\text{pp});$
 $(k, \text{sst}_A, \text{sst}_B, k) \leftarrow \text{AKE}(A(A.SK), O_A(O_A.SK),$
 $O_B(O_B.SK), B(B.SK))(\text{pp}, A.PK, B.PK, \text{APK});$
 τ_A is the transcript of the execution yielding sst_A from O_A 's point of view;
 τ_B is the transcript of the execution yielding sst_B from O_B 's point of view;

it holds that:

- $\{O_A.PK, O_B.PK, A.PK, B.PK\} \cup \{\Lambda_i.PK\}_{i=1}^n \subseteq \text{PPK},$
- for any oracle instance π_{O_A} such that $\pi_{O_A}.\tau = \tau_A, \pi_{O_A}.PID = (A, B), \pi_{O_A}.AID = (\Lambda_i)_{i=1}^n$, and $\pi_{O_A}.\text{sst} = \text{sst}$, we have: $\Pr[\text{Extract}(\pi_{O_A}, \text{PPK}) = k] = 1.$
- for any oracle instance π_{O_B} such that $\pi_{O_B}.\tau = \tau_B, \pi_{O_B}.PID = (A, B), \pi_{O_B}.AID = (\Lambda_i)_{i=1}^n$, and $\pi_{O_B}.\text{sst} = \text{sst}$, we have: $\Pr[\text{Extract}(\pi_{O_B}, \text{PPK}) = k] = 1.$

We stress the fact that our extractor is unbounded – indeed it must be so, or otherwise the property of key-security would be violated (the extractor would allow the operator to find the session key).

Definition 11. For any lawful interception key-exchange scheme LIKE that admits a key extractor Extract the advantage of an adversary \mathcal{A} in the honest-operator game $\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda)$ in Table 3 is defined as: $\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda) =$

$$\mathbb{P} \left[\begin{array}{l} (k_*, \pi_O, \text{PPK}) \leftarrow \text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda); \\ k \leftarrow \text{Extract}(\pi_O, \text{PPK}) \end{array} : \begin{array}{l} k \neq \perp \wedge k_* \neq \perp \\ \wedge k \neq k_* \end{array} \right].$$

A LIKE scheme is honest-operator secure if, for all adversaries running in time polynomial in λ , $\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda)$ is negligible as a function of λ .

5 Our protocol

To instantiate our LIKE scheme, we will need the following primitives (for which some background is given in Section 2):

- A signature scheme $\text{DS} = (\text{SGen}, \text{SSig}, \text{SVer})$;
- A non-interactive zero-knowledge proof of knowledge (NIPoK, NIPoKver) for a cyclic group $\mathbb{G}_1 = \langle g_1 \rangle$ allowing to prove knowledge of the discrete logarithm of a value $y = g_1^x$ for some private witness x . We denote this as $\text{NIPoK}\{x : y = g_1^x\}.$

- A signature of knowledge scheme (SoK, SoKver) allowing to prove the knowledge of a discrete logarithm in a second cyclic group $\mathbb{G}_2 = \langle g_2 \rangle$.
- A non-interactive zero-knowledge proof of knowledge (NIPoK, NIPoKver) for two groups $\mathbb{G}_1 = \langle g_1 \rangle$ and $\mathbb{G}_T = \langle g_T \rangle$ of same order p allowing to prove knowledge of the discrete logarithm of two values $y_1 = g_1^x$ and $y_T = g_T^x$ for some private witness x . We denote this by $\text{NIPoK} \{x : y_1 = g_1^x \wedge y_T = g_T^x\}$.

The proof and the signature of knowledge of a discrete logarithm can be instantiated with Schnorr's protocol [32]. This protocol consists in three passes: a commitment, a challenge and a response. The Fiat-Shamir heuristic allows to change Schnorr's protocol into a non-interactive proof by using the hash of the statement concatenated with the commitment as challenge, and into a signature of knowledge by adding the message in this hash [15]. The proof of the discrete logarithm equality can be instantiated with The Fiat-Shamir heuristic applied on the protocol of Chaum and Pedersen [16].

Our scheme follows the syntax in Section 3. We divide its presentation in four main components: (a) setup and key generation, (b) authenticated key-exchange, (c) public verification, and (d) lawful interception.

Setup and key generation. This part instantiates the four following algorithms of the LIKE syntax presented in Section 3.

- $\text{Setup}(1^\lambda)$: Based on λ , choose $\mathbb{G}_1 = \langle g_1 \rangle$, $\mathbb{G}_2 = \langle g_2 \rangle$, and \mathbb{G}_T , three groups of prime order p of length λ , $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T$ a type 2 bilinear mapping, and output $\text{pp} = (1^\lambda, \mathbb{G}_1, \mathbb{G}_2, \mathbb{G}_T, e, p, g_1, g_2)$.
- $\text{UKeyGen}(\text{pp})$: Run $(\text{U.PK}, \text{U.SK}) \leftarrow \text{SGen}(1^\lambda)$ and return $(\text{U.PK}, \text{U.SK})$.
- $\text{OKeyGen}(\text{pp})$: Run $(\text{O.PK}, \text{O.SK}) \leftarrow \text{SGen}(1^\lambda)$ and return $(\text{O.PK}, \text{O.SK})$.
- $\text{AKeyGen}(\text{pp})$: Pick $\Lambda.\text{SK} \xleftarrow{\$} \mathbb{Z}_p^*$, compute $\Lambda.\text{pk} \leftarrow g_1^{\Lambda.\text{SK}}$ and $\Lambda.\text{ni} \leftarrow \text{NIPoK} \{ \Lambda.\text{SK} : \Lambda.\text{pk} = g_1^{\Lambda.\text{SK}} \}$, set $\Lambda.\text{PK} \leftarrow (\Lambda.\text{pk}, \Lambda.\text{ni})$ and return $(\Lambda.\text{PK}, \Lambda.\text{SK})$.

The setup algorithm is run only once, to generate the groups required to instantiate our bilinear pairing. The users and operators generate signature keys, to be used with our digital-signature scheme DS in the authenticated key-exchange step. Authorities not only generate a private and public key-pair, but have to prove (in zero-knowledge) that they know the private key. This prevents attacks in which one or multiple authorities try to “annihilate” the contribution of one or more of the authorities and break key security. We explain the role of this proof in more detail after presenting our scheme. All key-generation algorithms are run only once per party.

Authenticated key exchange. Every time two users communicate, they run the the AKE component of our LIKE syntax, $\text{AKE}(\Lambda(\text{A.SK}), \text{O}_A(\text{O}_A.\text{SK}), \text{O}_B(\text{O}_B.\text{SK}), \text{B}(\text{B.SK}))$ ($\text{pp}, \text{A.PK}, \text{B.PK}, \text{APK}$).

Our protocol is described in Figure 1. For readability we abbreviate the picture and “merge” O_A and O_B . Notice particularly that the operators' roles are very similar to one another, and neither O_A , nor O_B will learn the keys established by the end points.

In Figure 1, all the parties involved in the protocol (namely Alice, Bob, and their respective operators) will verify the public keys of the n authorities and abort if their associated NIZK proofs of knowledge are not correct. Clearly, however, if one user (say Alice) has already performed this step for a given authority, it can, in practice, skip the verification of its zero-knowledge proof during a specific future session.

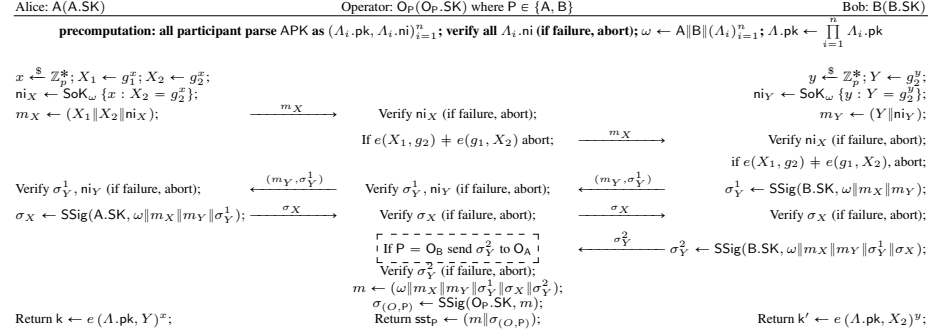


Fig. 1: The AKE component of our LIKE construction, with both operators O_A and O_B under a single heading (for readability). Unless specified, each operator runs in turn each line of the protocol, then forwards the message to the next participant or else aborts. The only message *not* forwarded by O_A to Alice is marked in the dashed box.

The heart of the protocol is the exchange between Alice and Bob. Alice generates a secret x and sends $X_1 = g_1^x, X_2 = g_2^x$ to Bob, with an associated signature of knowledge linking this key share to ω . Bob proceeds in the same manner, sampling a secret y and sending to Alice $Y = g_2^y$ and a signature of knowledge. The two parties also send to each other a signature over the transcript, thus authenticating each other and verifying the integrity of the conversation.

The two operators sit in between the two users. They verify the messages that they receive; if the message is correct, they forward it. If not, they abort the protocol. An exception is Bob's last message, which is verified by both operators, but does not reach Alice. The two operators check the signatures of knowledge, the digital signatures and the fact that the two values sent by Alice share the same discrete logarithm (by doing one pairing computation).

Given the public value Y , the public keys of the authorities, and her private exponent x , Alice computes her session key as a bilinear map applied on the product of all the public keys of the authorities and Bob's value Y , all raised to her secret x : $k = e(\prod_{i=1}^n A_i.pk, Y)^x$. Due to bilinearity, this is equal to $e(\prod_{i=1}^n A_i.pk, g_2^y)^x = e(\prod_{i=1}^n A_i.pk, g_2^x)^y = e(\prod_{i=1}^n A_i.pk, X_2)^y$, which is what Bob computes on his end. This indicates the correctness of our scheme with respect to the endpoints. The users compute the pairing on public values, and then exponentiate the result with a secret value. One can delegate the pairing to a faster, less secure environment (such as the phone), and do the exponentiation on the SIM card.

The operator's output is the session state sst , the signed session transcript from their point of view.

Verification. We instantiate Verify as follows:

- $\text{Verify}(pp, sst, A.PK, B.PK, O.PK, APK) \rightarrow b$: Parse APK as a set $(A_i.PK)_{i=1}^n$ and parse each $A_i.PK$ as $(A_i.pk, A_i.ni)$, set $\omega \leftarrow A \| B \| (A_i)_{i=1}^n$. Parse sst as $\omega' \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2 \| \sigma_O, m_X$ as $X_1 \| X_2 \| ni_X$ and m_Y as $Y \| ni_Y$. if:
 - For all $i \in [1, n]$, $\text{NIPoKver}(A_i.pk, A_i.ni) = 1$;

- $e(X_1, g_2) = e(g_1, X_2)$;
 - $\text{SoKver}(\omega, (g_2, X_2), \text{ni}_X) = 1$;
 - $\text{SoKver}(\omega, (g_2, Y), \text{ni}_Y) = 1$;
 - $\text{SVer}(\text{B.PK}, \sigma_Y^1, \omega \| m_X \| m_Y) = 1$;
 - $\text{SVer}(\text{A.PK}, \sigma_X, \omega \| m_X \| m_Y \| \sigma_Y^1) = 1$;
 - $\text{SVer}(\text{B.PK}, \sigma_Y^2, \omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X) = 1$;
 - $\text{SVer}(\text{O.PK}, \sigma_O, \omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2) = 1$;
- then the algorithm returns 1, else it returns 0.

Intuitively, one verifies that Alice and Bob participated in the session by verifying their signatures. The same can be said about the operator that output sst (which could be either O_A or O_B), whose final signature is verifiable with its public key. In addition, the verification algorithm performs the checks required by the protocol from the operator: verifying the proofs presented by the authorities, checking that X_1 and X_2 contain the same discrete logarithm ($e(X_1, g_2) = e(g_1, X_2)$), and verifying all the signatures of knowledge.

Lawful Interception. In Section 3, lawful interception for our protocol follows an extract-and-open strategy. We employ two algorithms, one run by each authority in order to generate a trapdoor, and the other run by one or a multitude of parties, cumulatively in possession of all the trapdoors for a given session.

- $\text{TDGen}(\text{pp}, \Lambda.\text{SK}, \text{sst})$: Parse sst as $(\omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2 \| \sigma_O)$. Compute $\Lambda.t_1 \leftarrow e(X_1, Y)^{\Lambda.\text{SK}}$, $\Lambda.t_2 \leftarrow \text{NIPoK} \left\{ \Lambda.\text{SK} : \Lambda.\text{PK} = g_1^{\Lambda.\text{SK}} \wedge \Lambda.t_1 = e(X_1, Y)^{\Lambda.\text{SK}} \right\}$ and $\Lambda.t \leftarrow (\Lambda.t_1, \Lambda.t_2)$, and returns $\Lambda.t$.
- $\text{Open}(\text{pp}, \text{sst}, \text{APK}, \mathcal{T})$: Parse \mathcal{T} as $(\Lambda_i.t)_{i=1}^n$, parse sst as $\text{A} \| \text{B} \| (\Lambda_i)_{i=1}^n \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2 \| \sigma_O$, m_X as $X_1 \| X_2 \| \text{ni}_X$ and m_Y as $Y \| \text{ni}_Y$, APK as $(\Lambda_i.\text{pk})_{i=1}^n$, parse each $\Lambda_i.\text{PK}$ as $(\Lambda_i.\text{pk}, \Lambda_i.\text{ni})$, each $\Lambda_i.t$ as $(\Lambda_i.t_1, \Lambda_i.t_2)$ and verify the non-interactive proof of knowledge: $\text{NIPoKver}((g_1, \Lambda_i.\text{pk}, (X_1, Y), \Lambda_i.t_1), \Lambda_i.t_2)$; if any verification fails, the Open algorithm returns \perp . Compute $k \leftarrow \prod_{i=1}^n (\Lambda_i.t_1)$ and return it.

Informally, each authority Λ_i generates as trapdoor the value $\Lambda_i.t_1 = e(X_1, Y)^{\Lambda_i.\text{SK}} = e(g_1, g_2^{xy})$. In order to recover the session key (by means of the Open algorithm), all the trapdoors from $\Lambda_1.t_1$ to $\Lambda_n.t_1$ are multiplied. We obtain $\prod_{i=1}^n \Lambda_i.t_1 = \prod_{i=1}^n e(g_1^x, g_2^y)^{\Lambda_i.\text{SK}} = e(g_1, g_2^{xy})^{\sum_{i=1}^n \Lambda_i.\text{SK}} = e(g_1^{\sum_{i=1}^n \Lambda_i.\text{SK}}, g_2^{xy}) = e(\prod_{i=1}^n \Lambda_i.\text{PK}, X_2)^y$. We now recognize the key as computed by Bob. Our scheme is correct with respect to LI.

Complexity. In our key exchange protocol, each party runs in linear time in the number of authorities. This complexity is due to the verification of the zero-knowledge proof for each authority public key, and the computation of the product of these public keys. We stress that the parties can pre-computed the proofs verifications and the product, moreover, while the parties run the protocol for the same authority set, they have not to recompute this, meaning that in practice, our protocol is in constant time most of the time. Finally, we note that the trapdoor generation algorithm also runs in constant time.

6 Proofs

We show that the LIKE scheme we propose in Section 5 has the non-frameability, key-security and honest operator properties defined in Section 4. We provide proof sketches in this section and full proofs in the appendix. However, we first provide an insight into the role of some of our less-obvious building blocks, namely the proofs and signatures of knowledge.

Insight: proofs and signatures of knowledge. During the protocol, Alice and Bob use the two signatures of knowledge on the public parameters of the session (denoted ω): $\text{SoK}_\omega \{x : X_2 = g_2^x\}$ and $\text{SoK}_\omega \{y : Y = g_2^y\}$. In the absence of these proofs, the adversary could recover the key of an honest session through the following attack. Let X_1, X_2, Y be the values honestly generated by the users in the targeted session. Choose two random values r and s , run the protocol using $X'_1 \leftarrow X_1^r, X'_2 \leftarrow X_2^r$ and $Y' \leftarrow Y^s$ for two corrupted users and the same set of authorities, obtain the corresponding sst, and run the RevealTD oracle on sst for each authority. Using the algorithm Open, the adversary obtains: $k' = \prod_{i=1}^n (X'_1, Y'_2)^{A_i.SK} = (\prod_{i=1}^n (X_1, Y_2)^{A_i.SK})^{r \cdot s}$. It can then compute the targeted key as $k = (k')^{1/r \cdot s}$. However, this attack cannot be done by the adversary if it must prove its knowledge of the discrete logarithm of X'_2 and Y' . Moreover, it cannot reuse X_2 and Y together with the signature of the honest user, since that signature uses the identity of the users as a message.

Each authority public key $A.pk$ must be coupled with a proof of knowledge $\text{NIPoK} \{A.SK : A.pk = g_1^{A.SK}\}$. Without this proof, an authority A_j in the set $(A_i)_{i=1}^n$ could pick a random r and compute: $A_j.pk = g_1^r / \left(\prod_{i=1; i \neq j}^n A_i.pk \right)$. In this case, the secret key computed by Alice will be: $k = e \left(\prod_{i=1}^n A_i.pk, Y \right)^x = e(g_1^r, g_2^y)^x = e(X_1, Y)^r$. Then the dishonest authority can discover the key by computing $e(X_1, Y)^r$. This attack is not possible if the authority must prove the knowledge of its secret key.

Finally, the proof $\text{NIPoK} \{A.SK : A.SK = g_1^{A.SK} \wedge A.t_1 = e(X_1, Y)^{A.SK}\}$, which is a part of the trapdoor generated by an authority, ensures that the trapdoor has been correctly formed. This is essential for the honest-operator property: without this proof, the authority can output a fake trapdoor, thus distorting the output of the algorithm Open.

Proofs. We proceed with the formal security statements and proofs. In the following, we define $\text{sid} := X_2 \| Y$ (see Figure 1).

We define the algorithm $\text{IdentifySession}(\text{sst}, \pi_P^j)$ for some party P and integer j as follows. Parsing sst as:

$$A \| B \| (A_i)_{i=1}^n \| X_1 \| X_2 \| \text{ni}_X \| Y \| \text{ni}_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2 \| \sigma_O,$$

the algorithm $\text{IdentifySession}(\text{sst}, \pi_P^j)$ returns 1 iff:

- $X_2 \| Y = \pi_P^j.\text{sid}$,
- if π_P^j plays the role of Alice then $P = A$ and $\pi_P^j.\text{PID} = B$, else $\pi_P^j.\text{PID} = A$ and $P = B$, and
- $\pi_P^j.AID = (A_i)_{i=1}^n$.

Theorem 1. *If our protocol is instantiated with an EUF-CMA signature scheme DS, and with extractable and zero know-ledge proofs/signature of knowledge, and if the BDDH assumption holds, then our protocol is key-secure. Moreover, for all PPT adversaries \mathcal{A} doing at most q_r queries to the oracle Register, q_{ns} queries to the oracle NewSession, q_s queries to the oracle Send and q_t queries to the oracle RevealTD, we have:*

$$\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda) \leq \frac{q_s^2}{p} + q_{ns} \cdot q_r^2 \cdot \left(\text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda) + q_{ns} \cdot q_r \cdot \left((2 \cdot q_t + q_s) \cdot \epsilon_{\text{SoK}}(\lambda) + q_r \cdot \epsilon_{\text{NIPoK}}(\lambda) + \text{Adv}^{\text{BDDH}}(\lambda) \right) \right).$$

Proof sketch. We begin by proving that the adversary has a negligible probability of winning the key-security experiment by querying the oracle Test on an instance that matches no other instance. Notably, if the tested instance does not abort the protocol, the adversary will have to break the EUF-CMA of the signature scheme to generate the expected signatures without using a matching session.

Thus, the targeted instance must have a matching one. By key-freshness, \mathcal{A} must test a key generated by two honest users, such that the trapdoor of at least one honest authority has never been queried to the oracle RevealTD. We prove (by a reduction) that \mathcal{A} can only win by breaking the BDDH assumption. Let $(W_*, X_*, Y_*, W'_*, X'_*, Y'_*, Z_*)$ be a BDDH instance. We set W_* as the part of the public key $\Lambda.\text{pk}$ of the honest authority, and we set X_2 as X'_* , X_1 as X_* and Y as Y'_* for the session that matches the tested instance. Then, we build the key as follows, where Λ is the honest authority: $k \leftarrow Z_* \prod_{i=1; \Lambda_i \neq \Lambda}^n e(X_*, Y'_*)^{\Lambda_i \cdot \text{SK}}$. To compute the secret keys of the authorities controlled by the adversary, we run the extractor on the proofs of knowledge of the discrete logarithm of the public keys $\Lambda_i.\text{PK}$. If Z_* is a random value, k will be random for the adversary, else $Z_* = e(X_*, Y'_*)^{\Lambda \cdot \text{SK}}$. Moreover, we simulate the oracle RevealTD on sessions with values X and Y chosen by the adversary by using the extractor on the signatures of knowledge of their discrete logarithms.

Proof. (Th.1: Key-security). We show that $\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda)$ is negligible for any PPT adversary \mathcal{A} , and thus our LIKE scheme is key-secure, using the following sequence of games:

Game G_0 : This game is the same as $\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda)$.

Game G_1 : This game is similar to G_0 , but aborts if the Send oracle returns twice the same element as X_2 or Y . An abort only happens if two out of the q_s queried instances choose the same randomness from \mathbb{G}_2 (which is of size p), yielding:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_0] - \mathbb{P}[\mathcal{A} \text{ wins } G_1]| \leq q_s^2/p.$$

Let $\pi_{P_*}^{i*}$ denote a tested instance. Excluding collisions for X_2 and Y implies that $\pi_{P_*}^{i*}$ now has at most one matching instance. Indeed, suppose two or more instances matching $\pi_{P_*}^{i*}$ exist. We parse $\pi_{P_*}^{i*}.\text{sid}$ as $Z_0 \| Z_1$ where Z_i (for $i \in \{0, 1\}$) was generated by $\pi_{P_*}^{i*}$.

By Def. 6, all instances matching $\pi_{P_*}^{i_*}$ must sample the same $Z_{1-i} \in \mathbb{G}_2$ – impossible after G_1 .

Game G_2 : Let P_i be the i -th party instantiated by Register. Game G_2 proceeds as G_1 except that it begins by choosing $(u, v, w) \xleftarrow{\$} [1, q_{ns}] \times [1, q_r]^2$. If \mathcal{A} returns (i_*, P_*, d_*) such that, given $P'_* \leftarrow \pi_{P_*}^{i_*} \cdot \text{PID}$, we have $i_* \neq u$ or $P_* \neq P_v$ or $P'_* \neq P_w$, then G_2 aborts, returning a random bit. The adversary increases its winning advantage by a factor equalling the probability of guessing correctly:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_1] - 1/2| \leq q_{ns} \cdot q_r^2 |\mathbb{P}[\mathcal{A} \text{ wins } G_2] - 1/2|.$$

Game G_3 : Let (i_*, P_*, d_*) be the adversary's test session that G_2 guessed. Let $P'_* \leftarrow \pi_{P_*}^{i_*} \cdot \text{PID}$. Game G_3 works as G_2 , except that, if there exists no $\pi_{P'_*}^k$ matching $\pi_{P_*}^{i_*}$, the experiment aborts and returns a random bit. For any adversary \mathcal{A} :

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_2] - \mathbb{P}[\mathcal{A} \text{ wins } G_3]| \leq \text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda).$$

Assume to the contrary that there exists an adversary \mathcal{A} that wins G_2 with probability $\epsilon_{\mathcal{A}}(\lambda)$ by returning a guess (i_*, P_*, d_*) such that, setting $P'_* \leftarrow \pi_{P_*}^{i_*} \cdot \text{PID}$, no $k \in \mathbb{N}$ exists such that $\pi_{P_*}^{i_*}$ and $\pi_{P'_*}^k$ match. Game G_2 demands $P'_* \leftarrow P_w$ (guessed by G_2); key-freshness (Def. 7) requires P_w to be uncorrupted and ending in an accepting state. We use \mathcal{A} to build a PPT adversary \mathcal{B} that breaks the EUF-CMA security of DS with non-negligible probability. \mathcal{B} receives the verification key $\hat{\text{PK}}$, initializes $\mathcal{L}_S \leftarrow \emptyset$, and faithfully simulates G_2 to \mathcal{A} , except for \mathcal{A} 's following queries:

Oracle Register($P, \text{role}, \text{PK}$): If $P = P_w$ with $P.\text{PK} = \perp$, then \mathcal{B} sets $P.\text{PK} \leftarrow \hat{\text{PK}}$.

Oracle Send(π_P^i, m): There are two particular cases: $P = P_w$ and $P = P_*$. If $P = P_w$, then \mathcal{B} queries its $\text{Sig}(\cdot)$ oracle to answer \mathcal{A} 's queries. Depending on the role of P_w and the protocol step, \mathcal{B} runs one of: $\sigma_Y^1 \leftarrow \text{Sig}(\omega \| m_X \| m_Y)$, $\sigma_X \leftarrow \text{Sig}(\omega \| m_X \| m_Y \| \sigma_Y^1)$, or $\sigma_Y^2 \leftarrow \text{Sig}(\omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X)$. Here, if P_w is the initiator, $\omega = P_w \| P_w.\text{PID} \| (\Lambda_l)_{l=1}^n$; else $\omega = P_w.\text{PID} \| P_w \| (\Lambda_l)_{l=1}^n$. Moreover, $m_X = X_1 \| X_2 \| \text{ni}_X$ and $m_Y = Y \| \text{ni}_Y$. The message/signature pairs are stored in \mathcal{L}_S . Since $\text{sid} = X_2 \| Y$, the elements X_2 and Y , the identities P_w and $\pi_{P_w}^i \cdot \text{PID}$, and the set of authorities $\pi_{P_w}^i \cdot \text{AID} = (\Lambda_l)_{l=1}^n$ are parts of the message signed in σ_X , σ_Y^1 and σ_Y^2 .

If $P = P_*$, $i = i_*$, and $\pi_P^i \cdot \text{PID} = P_w$, if $\text{SVer}(P_w.\text{PK}, \sigma_Y^2, M_Y) = 1$ for $M_Y \leftarrow \omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X$, and $(M_Y, \sigma_Y^2) \notin \mathcal{L}_S$, \mathcal{B} aborts the game and returns (M_Y, σ_Y^2) . Otherwise, if $\text{SVer}(P_w.\text{PK}, \sigma_X, M_X) = 1$ for $M_X \leftarrow \omega \| m_X \| m_Y \| \sigma_Y^1$ and $(M_X, \sigma_X) \notin \mathcal{L}_S$, \mathcal{B} aborts the game and returns (M_X, σ_X) .

Oracle Corrupt(P): If $P = P_w$, \mathcal{B} aborts (due to G_2).

\mathcal{B} wins if it sends its challenger a message/signature pair $(M, \sigma) \notin \mathcal{L}_S$ such that $\text{SVer}(\hat{\text{PK}}, \sigma, M) = 1$ with $\hat{\text{PK}} = P_w.\text{PK}$. We first argue that \mathcal{A} must query Send on input $P = P_*$, $i = i_*$, and $\pi_P^i \cdot \text{PID} = P_w$, on message $M_Y = \omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X$ such that $\text{SVer}(P_w.\text{PK}, \sigma_Y^2, M_Y) = 1$, or on message $M_X \leftarrow \omega \| m_X \| m_Y \| \sigma_Y^1$ such that $\text{SVer}(P_w.\text{PK}, \sigma_X, M_X) = 1$. Indeed, if \mathcal{A} does not, the (honest) target instance $\pi_{P_*}^{i_*}$ rejects.

Now we can assume that \mathcal{A} has queried Send either with σ_X or with σ_Y as above. We have two cases: the submitted message/signature pair is in \mathcal{L}_S , or it is not. If the

latter happens, clearly \mathcal{B} wins. Assume that the former happens, *i.e.*, the signature σ_Y^2 or σ_X are in \mathcal{L}_S (generated by \mathcal{B} 's oracle). We recall that by assumption \mathcal{A} 's challenge instance has no matching instance, *i.e.*, there exists no $\pi_{P_w}^j$ such that $\pi_{P_w}^j.\text{sid} \neq \perp$ or $\pi_{P_w}^j.\text{sid} = \pi_{P_*}^{i*}.\text{sid}$, and $\pi_{P_w}^j.\text{AID} = \pi_{P_*}^{i*}.\text{AID}$. However, if $\pi_{P_w}^j.\text{sid} \neq \pi_{P_*}^{i*}.\text{sid}$, or $\pi_{P_w}^j.\text{AID} \neq \pi_{P_*}^{i*}.\text{AID}$, then the signature is generated for the fresh message and \mathcal{A} would win.

Thus, $\text{Adv}_{\text{DS}, \mathcal{B}}^{\text{EUF-CMA}}(\lambda) = \epsilon_{\mathcal{A}}(\lambda)$, concluding the proof.

After G_2 , G_3 , either a unique instance $\pi_{P_w}^r$ exists, matching $\pi_{P_*}^{i*}$ or the experiment returns a random bit.

Game G_4 : Game G_4 runs as G_3 except that it begins by picking $r \xleftarrow{\$} [1, q_{ns}]$ (a guess for the matching instance). If \mathcal{A} returns (i_*, P_*, d_*) such that $\pi_{P_*}^{i*}$ and $\pi_{P_w}^r$ do not match, then the experiment returns a random bit. The advantage of \mathcal{A} on G_4 increases w.r.t. that in G_3 by a factor equalling the correct guessing probability :

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_3] - 1/2| \leq q_{ns} |\mathbb{P}[\mathcal{A} \text{ wins } G_4] - 1/2|.$$

Game G_5 : Game G_5 proceeds as G_4 , except that it begins by picking $l \xleftarrow{\$} [1, q_r]$. If the l -th party queried to the oracle Register is not authority, or, if it is an authority (we will denote it A_l), and is either corrupted, or RevealTD is called on the query $(\text{sst}, A, B, (A'_i)_{i=1}^{n'}, l')$ such that $A_l = A'_l$ and $\text{IdentifySession}(\text{sst}, \pi_{P_v}^u.\text{sid}) = 1$, then the experiment aborts by returning a random bit. By key-freshness (Def. 7), if no index l exists such that A_l is uncorrupted ($A_l.\gamma = 0$) and RevealTD has never been called on the query $(\text{sst}, A, B, (A'_i)_{i=1}^{n'}, l')$ such that $A_l = A'_l$ and $\text{IdentifySession}(\text{sst}, \pi_{P_v}^j.\text{sid}) = 1$, then the experiment returns a random bit. Thus, the advantage of \mathcal{A} in G_5 is superior to that in G_4 by a factor equalling the guessing probability:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_4] - 1/2| \leq q_{ns} |\mathbb{P}[\mathcal{A} \text{ wins } G_5] - 1/2|.$$

Game G_6 : Let Ext denote the knowledge extractor of the signature of knowledge. This game is the same as G_5 except that it begins by initializing $\mathcal{L}[\] \leftarrow \emptyset$ and:

- each time the Send oracle generates an element $d \xleftarrow{\$} \mathbb{Z}_p^*$ and $D \leftarrow g_2^d$ together with a signature of knowledge $\text{ni}_D \leftarrow \text{SoK}_{\omega} \{d : D = g_2^d\}$, $\mathcal{L}[(D, \omega, \sigma_D)] \leftarrow d$;
- each time the oracles Send or RevealTD verify a valid signature of knowledge $\text{SoKver}(\omega, (g_2, D), \text{ni}_D) = 1$ in a query sending by \mathcal{A} with $\mathcal{L}[(D, \omega, \sigma_D)] = \perp$, it runs the key extractor Ext(λ) on \mathcal{A} in order to extract the witness d that matches the proof ni_D . If $g_2^d \neq D$ then the experiment aborts by returning a random bit, else it sets $\mathcal{L}[(D, \omega, \sigma_D)] \leftarrow d$.

The difference between G_5 and G_6 is the possibility of the extractor failing when it is called. Since RevealTD requires 2 calls (for the verification of sst) and Send, one at each query:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_5] - \mathbb{P}[\mathcal{A} \text{ wins } G_6]| \leq (2 \cdot q_t + q_s) \cdot \epsilon_{\text{SoK}}(\lambda).$$

Game G_7 : Let Ext denote the extractor of the NIZK proof of knowledge $\text{NIPoK} \{d : D = g_1^d\}$. Game G_7 runs as G_6 , except it begins by initializing an empty list $\mathcal{L}'[\] \leftarrow \emptyset$ and:

- Honest authority: if Register generates $(\Lambda.PK, \Lambda.SK)$ for an authority Λ , it sets $\mathcal{L}'[\Lambda.PK] \leftarrow \Lambda.SK$;
- Malicious authority: if Register receives a query $(\Lambda, \text{role}, PK)$ with $\text{role} = \text{authority}$ and $PK \neq \perp$, it sets $PK \leftarrow \Lambda.PK$ and parses PK as $(\Lambda.pk, \Lambda.ni)$. If $\text{NIPoKver}((g_1, \Lambda.pk), \Lambda.ni) = 1$, G_7 runs the extractor $\text{Ext}(\lambda)$ on \mathcal{A} to get the witness $\Lambda.SK$ for $\Lambda.ni$. If $g_1^{\Lambda.SK} \neq \Lambda.pk$ then the experiment aborts by returning a random bit, else it sets $\mathcal{L}'[\Lambda.PK] \leftarrow \Lambda.SK$.

Once more, the difference between the games is the possibility that Ext fails in at least one of the calls to the registration oracle, yielding:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_6] - \mathbb{P}[\mathcal{A} \text{ wins } G_7]| \leq q_r \cdot \epsilon_{\text{NIPoK}}(\lambda).$$

Finally, we claim that:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_7] - 1/2| \leq \text{Adv}_{\mathcal{A}}^{\text{BDDH}}(\lambda).$$

We prove this claim by reduction. Assume that \mathcal{A} wins G_7 with non-negligible probability $\epsilon_{\mathcal{A}}(\lambda)$. We show how to build an algorithm \mathcal{B} that breaks the BDDH problem with probability $\epsilon_{\mathcal{A}}(\lambda)$.

In what follows, $\text{Sim}_{\text{NIPoK}}$ denotes the simulator of the proofs of knowledge and Sim_{SoK} is the simulator of the signature of knowledge. For readability, if the context is clear, we use the same notation for the simulators of the two proof of knowledge systems we use to instantiate our scheme.

\mathcal{B} plays the BDDH with $(\mathbb{G}_1 = \langle g_1 \rangle, \mathbb{G}_2 = \langle g_2 \rangle, \mathbb{G}_T, e, p)$ and receives $(\hat{W}, \hat{X}, \hat{Y}, \hat{W}', \hat{X}', \hat{Y}', \hat{Z})$ from its challenger. Then \mathcal{B} sets $\text{pp} \leftarrow (1^\lambda, \mathbb{G}_1, \mathbb{G}_2, \mathbb{G}_T, e, p, g_1, g_2)$, it runs $\mathcal{A}(\text{pp})$ and simulates G_7 to \mathcal{A} as in the real game except for following cases.

Oracle Register(P, role, PK) $\rightarrow P.PK$: On the l -th query, if $\text{role} \neq \text{authority}$ or $PK \neq \perp$, \mathcal{B} aborts and returns a random bit (this is a faithful simulation since G_5), else it sets $\Lambda_l \leftarrow P$; $\Lambda_l.pk \leftarrow \hat{W}$; $\Lambda_l.ni \leftarrow \text{Sim}_{\text{NIPoK}}(g_1, \hat{W})$; $\Lambda_l.PK \leftarrow (\Lambda_l.pk, \Lambda_l.ni)$ and returns $\Lambda_l.PK$.

Oracle Send(π_P^i, m): If $P = P_*$ and $i = i_*$ (challenge instance, guessed in G_3), or if $P = P_w$ and $i = r$ (unique matching instance, guessed in G_4), then we distinguish two cases:

- P plays the role of Alice: then \mathcal{B} proceeds as in G_7 except for generating (X_1, X_2, ni_X) by setting $X_2 \leftarrow \hat{X}'$, $X_1 \leftarrow \hat{X}$ and by running $\text{ni}_X \leftarrow \text{Sim}_{\text{SoK}}(\omega, (g_2, \hat{X}'))$, where $\omega = (P_v \| P_w \| \pi_{P_v}^i.AID)$.
- P plays the role of Bob: then \mathcal{B} proceeds as in G_7 except for generating (Y, ni_Y) by setting $Y \leftarrow \hat{Y}'$ and by running $\text{ni}_Y \leftarrow \text{Sim}_{\text{SoK}}(\omega, (g_2, \hat{Y}'))$, where $\omega = (P_w \| P_v \| \pi_{P_w}^i.AID)$.

At the end of the protocol, the oracle does not compute k and sets $\pi_{P_*}^i.k \leftarrow \perp$.

Oracle Test(π_P^i): If $P = P_*$ and $i = i_*$, parsing $\pi_P^i.AID$ as $(A'_j)_{j=1}^n$, \mathcal{B} sets $\text{AuthSet} \leftarrow \{A'_j\}_{j=1}^n \setminus \{\Lambda_l\}$, it computes: $\tilde{k} \leftarrow \left(\prod_{\Lambda \in \text{AuthSet}} e(\hat{X}, \hat{Y}')^{\mathcal{L}'[\Lambda.PK]} \right) \cdot \hat{Z}$, and returns it.

Note that if \mathcal{B} 's challenge bit is 0, then $\hat{Z} = e(\hat{X}, \hat{Y}')^{\Lambda_l.SK}$ for $\Lambda_l.SK$ such that $g_1^{\Lambda_l.SK} = \Lambda_l.pk$, so \tilde{k} is the real key expected for $\pi_{P_*}^{i_*}$. If, however, the challenge bit of \mathcal{B} is 1, then \hat{Z} is a random value.

Oracle RevealTD($\text{sst}, A, B, O, (A'_i)_{i=1}^n, l'$):

- if $\text{IdentifySession}(\text{sst}, \pi_{P_*}^{i*}) = 1$ and $\Lambda_{l'} = \Lambda_l$ (challenge instance, honest authority), then \mathcal{B} aborts returning a random bit, as for key freshness;
- if $\text{IdentifySession}(\text{sst}, \pi_{P_*}^{i*}) = 1$ and $\Lambda_{l'} \neq \Lambda_l$ (challenge instance, other authority), then \mathcal{B} knows the secret key of $\Lambda_{l'}$. It acts as in G_7 except that it computes $\Lambda_{l'}.t_1 \leftarrow e(\hat{X}, \hat{Y}')^{\mathcal{L}'[\Lambda_{l'}.PK]}$, $\Lambda_{l'}.t_2 \leftarrow \text{Sim}_{\text{NIPoK}}(g_1, \Lambda_{l'}.pk, g_T, \Lambda_{l'}.t_1)$ and $\Lambda_{l'}.t \leftarrow (\Lambda_{l'}.t_1, \Lambda_{l'}.t_2)$;
- if $\text{IdentifySession}(\text{sst}, \pi_{P_*}^{i*}) \neq 1$ and $\Lambda_{l'} = \Lambda_l$ (non-challenge session, honest authority), then \mathcal{B} parses sst as $(\omega' \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2 \| \sigma_O)$, m_X as $X_1 \| X_2 \| \text{ni}_X$ and m_Y as $Y \| \text{ni}_Y$, and sets $\omega \leftarrow A \| B \| (\Lambda_i')_{i=1}^n$. Note that here, $\text{IdentifySession}(\text{sst}, \pi_{P_v}^u) \neq 1$ implies that $X_2 \| Y \neq \hat{X}' \| \hat{Y}'$, or $A \| B \neq P_* \| P_w$ (or $A \| B \neq P_w \| P_*$ depending on who plays the role of Alice and Bob), or $(\Lambda_i')_{i=1}^n \neq \pi_{P_*}^{i*} \cdot \text{AID}$. It acts as in G_7 except that:
 - if $A \| B \| (\Lambda_i')_{i=1}^n \neq P_* \| P_w \| \pi_{P_*}^{i*} \cdot \text{AID}$ (or $P_w \| P_* \| \pi_{P_w}^r \cdot \text{AID}$, if P_* plays the role of Bob), then \mathcal{B} computes $\Lambda_l.t_1 \leftarrow e(\Lambda_l.pk, X)^{\mathcal{L}[(Y, \omega, \sigma_Y)]}$; $\Lambda_l.t_2 \leftarrow \text{Sim}_{\text{NIPoK}}(g_1, \Lambda_l.pk, g_T, \Lambda_l.t_1)$ and $\Lambda_l.t \leftarrow (\Lambda_l.t_1, \Lambda_l.t_2)$. Here, $\mathcal{L}[(Y, \omega, \sigma_Y)]$ is always defined since $\omega \neq P_* \| P_w \| \pi_{P_*}^{i*} \cdot \text{AID}$ (or $P_w \| P_* \| \pi_{P_w}^r \cdot \text{AID}$ if P_* plays the role of Bob).
 - if $A \| B \| (\Lambda_i')_{i=1}^n = P_* \| P_w \| \pi_{P_*}^{i*} \cdot \text{AID}$ (or $P_w \| P_* \| \pi_{P_w}^r \cdot \text{AID}$, if P_* plays the role of Bob) and $X_2 = \hat{X}'$, then $Y \neq \hat{Y}'$, and \mathcal{B} computes $\Lambda_l.t_1 \leftarrow e(\Lambda_l.pk, \hat{X}')^{\mathcal{L}[(Y, \omega, \sigma_Y)]}$; $\Lambda_l.t_2 \leftarrow \text{Sim}_{\text{NIPoK}}(g_1, \Lambda_l.pk, g_T, \Lambda_l.t_1)$ and $\Lambda_l.t \leftarrow (\Lambda_l.t_1, \Lambda_l.t_2)$. In this case, $\mathcal{L}[(Y, \omega, \sigma_Y)]$ is always defined because $Y \neq \hat{Y}'$.
 - else if $A \| B \| (\Lambda_i')_{i=1}^n = P_* \| P_w \| \pi_{P_*}^{i*} \cdot \text{AID}$ (or $P_w \| P_* \| \pi_{P_w}^r \cdot \text{AID}$, if P_* plays the role of Bob) and $Y = \hat{Y}'$, $X_2 \neq \hat{X}'$. \mathcal{B} sets $\Lambda_{l'}.t_1 \leftarrow e(\Lambda_l.pk, \hat{Y}')^{\mathcal{L}[(X_2, \omega, \sigma_X)]}$; $\Lambda_{l'}.t_2 \leftarrow \text{Sim}_{\text{NIPoK}}(g_1, \Lambda_l.pk, g_T, \Lambda_{l'}.t_1)$ and $\Lambda_{l'}.t \leftarrow (\Lambda_{l'}.t_1, \Lambda_{l'}.t_2)$. In this case, $\mathcal{L}[(X_2, \omega, \sigma_X)]$ is always defined because $X_2 \neq \hat{X}'$.

At the end of the game, \mathcal{B} forwards the bit b_* received from \mathcal{A} . The game is perfectly simulated for \mathcal{A} , and \mathcal{B} wins its BDDH challenge with the same probability as \mathcal{A} wins G_7 , which concludes the claim. Finally, by composing the probability of all the games, we obtain the bound of $\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{KS}}(\lambda)$.

Theorem 2. *If our protocol is instantiated with an EUF-CMA signature scheme DS, then our protocol is non-frameable. Moreover, for all PPT adversaries \mathcal{A} , doing at most q_r queries to the oracle Register, we have:*

$$\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda) \leq q_r \cdot \text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda).$$

Proof sketch. To win the non-frameability experiment, the adversary has to build a valid session state sst for a given user, containing a valid signature of this user. We prove this theorem by reduction: assuming that an adversary is able to break the non-frameability, since this adversary generates a valid signature for a user, we can use it to break the EUF-CMA security.

Proof. (Th2. Non-frameability).

Game G_0 : The original game $\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{NF}}(\lambda)$.

Game G_1 : Let P_i the i -th party instantiated by Register. Game G_1 runs as G_0 except that it begins by picking $u \xleftarrow{\$} [1, q_r]$. If \mathcal{A} returns (sst, P) such that $P \neq P_u$, then G_1

aborts and returns 0. The advantage of \mathcal{A} on G_1 increases w.r.t. that in G_0 by a factor equal to the probability of correctly guessing u :

$$\mathbb{P}[\mathcal{A} \text{ wins } G_0] \leq q_r \cdot \mathbb{P}[\mathcal{A} \text{ wins } G_1].$$

We prove that $\mathbb{P}[\mathcal{A} \text{ wins } G_1] \leq \text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda)$ by reduction. Assume there exists an adversary \mathcal{A} that wins G_1 with probability $\epsilon_{\mathcal{A}}(\lambda)$ by returning (sst_*, P_*) such that $\exists (A, B) \in \text{USERS}^2, O \in \text{OPS}, n \in \mathbb{N}$ and $(\Lambda_i)_{i=1}^n \in \text{AUTH}^n$ such that:

- $P_* = P_u; P_* \in \{A, B\}; P_*. \gamma = 0;$
- $\text{Verify}(\text{pp}, \text{sst}_*, A.\text{PK}, B.\text{PK}, O.\text{PK}, (\Lambda_i.\text{PK})_{i=1}^n) = 1;$
- $\forall i, \text{if } \pi_{P_*}^i \neq \perp \text{ then } \text{IdentifySession}(\text{sst}_*, \pi_{P_*}^i) = 0 \text{ or } \pi_{P_*}^i.\alpha = 0.$

We build a PPT adversary \mathcal{B} breaking the EUF-CMA security of DS with non-negligible probability. \mathcal{B} receives the verification key $\hat{\text{PK}}$ from its challenger, initializes a set $\mathcal{L}_S \leftarrow \emptyset$, then it simulates G_1 to \mathcal{A} as in the real game, except in the following situations:

Oracle Register($P, \text{role}, \text{PK}$): If $P = P_*$ (as guessed before) and $\text{PK} = \perp$, \mathcal{B} sets $P_*. \text{PK} \leftarrow \hat{\text{PK}}$.

Oracle Send($\pi_{P_*}^i, m$): If $P = P_*$, \mathcal{B} simulates the oracle faithfully except it uses its $\text{Sig}(\cdot)$ oracle to produce signatures. Depending on the role of P_* and the protocol step, it runs either $\sigma_Y^1 \leftarrow \text{Sig}(\omega \| m_X \| m_Y)$, $\sigma_X \leftarrow \text{Sig}(\omega \| m_X \| m_Y \| \sigma_Y^1)$ or $\sigma_Y^2 \leftarrow \text{Sig}(\omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X)$, where $\omega = A \| B \| (\Lambda_i)_{i=1}^n$, $m_X = X_1 \| X_2 \| \text{ni}_X$ and $m_Y = Y \| \text{ni}_Y$. The message/signature pairs are stored in \mathcal{L}_S .

Since $\text{sid} = X_2 \| Y$, then $X_2, Y, A = P_*, B = \pi_{P_*}^i.\text{PID}$ (or $A = \pi_{P_*}^i.\text{PID}, B = P_*$, depending on who plays the role of Alice and Bob), and $(\Lambda_i)_{i=1}^n = \pi_{P_*}^i.\text{AID}$ are parts of the messages signed in σ_X, σ_Y^1 , and σ_Y^2 .

Oracle Corrupt(P): If $P = P_*$, then \mathcal{B} aborts.

We parse sst_* as $\omega' \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2 \| \sigma_O$, m_X as $X_1 \| X_2 \| \text{ni}_X$ and m_Y as $Y \| \text{ni}_Y$. For readability, we denote by M_X the value $\omega' \| X_1 \| X_2 \| \text{ni}_X \| Y \| \text{ni}_Y \| \sigma_Y^1$, and M_Y the value $M_X \| \sigma_X$. If $\text{Verify}(\text{pp}, \text{sst}_*, A.\text{PK}, B.\text{PK}, O.\text{PK}, (\Lambda_i.\text{PK})_{i=1}^n) = 1$, we have that $\text{SVer}(A.\text{PK}, \sigma_X, M_X) = 1$ and $\text{SVer}(B.\text{PK}, \sigma_Y^2, M_Y) = 1$.

If $\forall i, \pi_{P_*}^i \neq \perp$, then $\text{IdentifySession}(\text{sst}_*, \pi_{P_*}^i) = 0$ or $\pi_{P_*}^i.\alpha = 0$. Thus the oracle Send never outputs valid σ_X or σ_Y^2 by using the $\text{Sig}(\cdot)$ oracle on messages matching the session identifier $\text{sid} = X_2 \| Y$, the identities A, B , and $(\Lambda_i)_{i=1}^n$ together. We have two cases. If $P_* = A$, then $(M_X, \sigma_X) \notin \mathcal{L}_S$. If $P_* = B$, then $(M_Y, \sigma_Y^2) \notin \mathcal{L}_S$. Finally, if $P_* = A$, then \mathcal{B} returns (M_X, σ_X) . If $P_* = B$, then \mathcal{B} returns (M_Y, σ_Y^2) .

The experiment is perfectly simulated for \mathcal{A} , and if \mathcal{A} wins, then \mathcal{B} returns a fresh and valid signature. Hence, we get $\text{Adv}_{\text{DS}, \mathcal{B}}^{\text{EUF-CMA}}(\lambda) = \epsilon_{\mathcal{A}}(\lambda)$, which concludes the proof.

Theorem 3. *If our protocol is instantiated with an EUF-CMA signature scheme DS and with extractable and zero-knowledge proofs/signature of knowledge, then our protocol is honest-operator. Moreover, for all PPT adversaries \mathcal{A} doing at most q_r queries to the oracle Register, we have:*

$$\text{Adv}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda) \leq q_r \cdot \epsilon_{\text{NIPoK}}(\lambda) + \text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda).$$

Proof sketch. The first step of the HO proof is to design a key extractor, which takes in input a session state sst , brute-forces the discrete logarithm of Bob's Y , then computes the key as Bob would: $k = e(\prod_{i=1}^n \Lambda_i.\text{pk}, X_2)^y$. Our goal is to prove that this is the key the authorities would retrieve.

We first show (by reduction) that the adversary can only build by itself a valid sst (that may match a fake authority set) with negligible probability. Namely, if an adversary can output valid signatures for an honest operator, then we can use it to break the EUF-CMA of the signature scheme.

Moreover, for any authority Λ and any values X_1 and Y , the proof of knowledge of a trapdoor ensures that $g_1^{\Lambda.\text{SK}} = \Lambda.\text{pk}$ and $\Lambda.t_1 = e(X_1, Y)^{\Lambda.\text{SK}}$, which implies that $\Lambda.t_1 = e(\Lambda.\text{pk}, X_2)^y$ and: $k_* = \prod_{i=1}^n \Lambda_i.t_1 = e(\prod_{i=1}^n \Lambda_i.\text{pk}, X_2)^y$. Thus, to win the HO experiment (and return a key such that $k \neq k_*$), the adversary must produce a proof on a false statement, which happens with negligible probability.

Proof. (Th3. Honest operator). For this proof, we first describe the (deterministic, unbounded) key-extractor algorithm $\text{Extract}(\pi_O, \text{PPK})$ which, given as input an operator instance π_O and a set of public keys, outputs the same key k as the (honest) instances of Alice and Bob.

$\text{Extract}(\pi_O, \text{PPK})$ works as follows:

- Parse $\pi_O.\tau$ as $(m_X, m'_X, (m_Y, \sigma_Y^1), (m'_Y, \sigma_Y^1), \sigma_X, \sigma'_X, \sigma_Y^2, \sigma_O)$ if O is the operator of Alice, $(m_X, m'_X, (m_Y, \sigma_Y^1), (m'_Y, \sigma_Y^1), \sigma_X, \sigma'_X, \sigma_Y^2, \sigma_O)$ otherwise, m_X as $(X_1 \| X_2 \| \text{ni}_X)$, m_Y as $(Y \| \text{ni}_Y)$, $\pi_O.\text{PID}$ as (A, B) and $\pi_O.\text{AID}$ as $(\Lambda_i)_{i=1}^n$,
- Set $y = 0$. While $g_2^y \neq Y$, do $y = y + 1$. Output y (which exists, since $Y \in \mathbb{G}_2$).
- For all $i \in \llbracket 1, n \rrbracket$, if $\Lambda_i.\text{PK} \notin \text{PPK}$ then abort, else parse each $\Lambda_i.\text{PK}$ as $(\Lambda_i.\text{pk}, \Lambda_i.\text{ni})$.
- Compute $k \leftarrow e(\prod_{i=1}^n \Lambda_i.\text{pk}, X_2)^y$ and return k .

This extractor is correct because k is computed exactly as for Bob in the real protocol.

Informally, in the HO security game, the adversary aims to output a session state sst and a serie of trapdoors such that (i) Verify accepts the session state and (ii) when applied to sst and the trapdoors, Open returns a key matching the one extracted from Extract on the corresponding operator instance. The adversary can corrupt all users and all authorities, but not the operator. We use the following sequence of games:

Game G_0 : The original game $\text{Exp}_{\text{LIKE}, \mathcal{A}}^{\text{HO}}(\lambda)$:

Game G_1 : Let O_i be the i -th oracle party output by Register. Game G_1 runs as G_0 except that it begins by picking $u \xleftarrow{\$} \llbracket 1, q_r \rrbracket$. If \mathcal{A} returns $(j_*, \text{sst}_*, A_*, B_*, O_*, (\Lambda_{*,i}, \Lambda_{*,i}.t)_{i=1}^n)$ such that $O_u \neq O_*$, then the experiment returns 0. The advantage of \mathcal{A} on G_1 will be minored by the advantage of G_0 multiplied by the probability of guessing correctly the operator O_* : $|\mathbb{P}[\mathcal{A} \text{ wins } G_0] - 1/2| \leq q_r |\mathbb{P}[\mathcal{A} \text{ wins } G_1] - 1/2|$.

Game G_2 : Let $(j_*, \text{sst}_*, A_*, B_*, O_*, (\Lambda_{*,i}, \Lambda_{*,i}.t)_{i=1}^n)$ be the guess of the adversary.

Game G_2 runs as G_1 , except that if for all $k \in \mathbb{N}$, $\pi_{O_*}^k.\text{sid} \neq \text{sid}_*$ or $\pi_{O_*}^k.\text{PID} \neq (A_*, B_*)$ or $\pi_{O_*}^k.\text{AID} \neq (\Lambda_{*,i})_{i=1}^n$, then G_2 aborts and returns 0. For any adversary \mathcal{A} ,

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_0] - \mathbb{P}[\mathcal{A} \text{ wins } G_1]| \leq \text{Adv}_{\text{DS}}^{\text{EUF-CMA}}(\lambda).$$

We prove this claim by reduction. Assume that there exists \mathcal{A} winning G_1 with probability $\epsilon_{\mathcal{A}}(\lambda)$ by returning a guess $(j_*, \text{sst}_*, A_*, B_*, O_*, (\Lambda_{*,i}, \Lambda_{*,i}.t)_{i=1}^n)$ such that for

all $k \in \mathbb{N}$, $\pi_{O_*}^k.\text{sid} \neq \text{sid}_*$ or $\pi_{O_*}^k.\text{PID} \neq (A_*, B_*)$ or $\pi_{O_*}^k.\text{AID} \neq (\Lambda_{*,i})_{i=1}^n$. We recall that if \mathcal{A} wins G_0 , then $O_*. \gamma = 0$ and $\text{Verify}(\text{pp}, \text{sst}_*, A_*. \text{PK}, B_*. \text{PK}, O_*. \text{PK}, (\Lambda_{*,i}. \text{PK})_{i=1}^n) = 1$.

Then, according to the rules of G_1 , if $O_* \neq O_u$, the experiment returns 0. We build a PPT adversary \mathcal{B} that breaks the EUF-CMA security of DS with non-negligible advantage. \mathcal{B} receives the verification key $\hat{\text{PK}}$, initializes an empty set $\mathcal{L}_S \leftarrow \emptyset$, and then simulates G_0 to $\mathcal{A}(\text{pp})$ faithfully, except in the following situations:

Oracle Register(P, role, PK): If the u -th registration query has $\text{role} \neq \text{operator}$, or $\text{PK} \neq \perp$ then \mathcal{B} aborts (as required since G_1). Else \mathcal{B} sets $O_u \leftarrow P$ and sets $O_u. \text{PK} \leftarrow \hat{\text{PK}}$.

Oracle Send(π_P^i, m): If $P = O_u$, then \mathcal{B} simulates this oracle faithfully, except that it queries its $\text{Sig}(\cdot)$ oracle to produce signatures, storing the message/signature pairs in \mathcal{L}_S . \mathcal{B} sets $M_O \leftarrow \omega \| m_X \| m_Y \| \sigma_Y^1 \| \sigma_X \| \sigma_Y^2$, runs $\sigma_O \leftarrow \text{Sig}(M_O)$, stores (M_O, σ_O) in \mathcal{L}_S , and sets $\text{sst} \leftarrow (M_O \| \sigma_O)$.

Oracle Corrupt(P): If $P = O_u$, then \mathcal{B} aborts.

Finally, \mathcal{A} returns $(j_*, \text{sst}_*, A_*, B_*, O_*, (\Lambda_{*,i}, \Lambda_{*,i}.t)_{i=1}^n)$, \mathcal{B} parses sst_* as $(\omega_* \| m_{*,X} \| m_{*,Y} \| \sigma_{*,Y}^1 \| \sigma_{*,X} \| \sigma_{*,Y}^2 \| \sigma_{*,O})$ and returns $(m_*, \sigma_*) = (\omega_* \| m_{*,X} \| m_{*,Y} \| \sigma_{*,Y}^1 \| \sigma_{*,X} \| \sigma_{*,Y}^2 \| \sigma_{*,O})$.

Note that G_1 is perfectly simulated for \mathcal{A} by \mathcal{B} . If for all $k \in \mathbb{N}$, $\pi_{O_*}^k.\text{sid} \neq \text{sid}_*$ or $\pi_{O_*}^k.\text{PID} \neq (A_*, B_*)$ or $\pi_{O_*}^k.\text{AID} \neq (\Lambda_{*,i})_{i=1}^n$, then sst_* was not returned by Send on an oracle instance, so $(m_*, \sigma_*) \notin \mathcal{L}_S$. Moreover, if \mathcal{A} wins G_1 , $\text{Verify}(\text{pp}, \text{sst}_*, A_*. \text{PK}, B_*. \text{PK}, O_*. \text{PK}, (\Lambda_{*,i}. \text{PK})_{i=1}^n) = 1$, implying that $\text{SVer}(O_*. \text{PK}, \omega_* \| m_{*,X} \| m_{*,Y} \| \sigma_{*,Y}^1 \| \sigma_{*,X} \| \sigma_{*,Y}^2 \| \sigma_{*,O}) = 1$, so σ_* is a valid signature on m_* for the key $\hat{\text{PK}}$. We also have that $O_* = O_u$.

If \mathcal{A} wins G_0 such that for all $k \in \mathbb{N}$, $\pi_{O_*}^k.\text{sid} \neq \text{sid}_*$ or $\pi_{O_*}^k.\text{PID} \neq (A_*, B_*)$ or $\pi_{O_*}^k.\text{AID} \neq (\Lambda_{*,i})_{i=1}^n$, then \mathcal{B} wins his game. Hence, $\text{Adv}_{\text{DS}, \mathcal{B}}^{\text{EUF-CMA}}(\lambda) = \epsilon_{\mathcal{A}}(\lambda)$, concluding the proof.

Game G_3 : In this game, Ext denotes the knowledge extractor of NIPoK $\{d : D_1 = g_1^d \wedge D_2 = g_T^d\}$.

Let $(j_*, \text{sst}_*, A_*, B_*, O_*,$

$(\Lambda_{*,i}, \Lambda_{*,i}.t)_{i=1}^n)$ be the guess of the adversary. Game G_3 runs as G_2 except that:

- G_3 begins by initializing an empty list $\mathcal{L}[\] \leftarrow \emptyset$.
- Each time it runs $\text{RevealTD}(\text{sst}, A, B, O, (\Lambda_i)_{i=1}^n, l) \rightarrow \Lambda_l.t$ to the adversary, G_3 sets $\mathcal{L}[\Lambda_l.t] \leftarrow \Lambda_l.\text{SK}$.
- After the guess of the adversary, it sets $\omega_* \leftarrow A_* \| B_* \| (\Lambda_{*,i})_{i=1}^n$, it parses the session state sst_* as $(\omega'_* \| m_{*,X} \| m_{*,Y} \| \sigma_{*,Y}^1 \| \sigma_{*,X} \| \sigma_{*,Y}^2 \| \sigma_{*,O})$, $m_{*,X}$ as $X_{*,1} \| X_{*,2} \| \text{ni}_{*,X}$ and $m_{*,Y}$ as $Y_* \| \text{ni}_{*,Y}$.
- For each $i \in [1, n]$ such that $\mathcal{L}[\Lambda_{*,i}.t] = \perp$, it parses $\Lambda_{*,i}. \text{PK}$ as $(\Lambda_{*,i}. \text{pk}, \Lambda_{*,i}. \text{ni})$ and $\Lambda_{*,i}.t$ as $(\Lambda_{*,i}.t_1, \Lambda_{*,i}.t_2)$. If $1 \leftarrow \text{NIPoKver}((g_1, \Lambda_{*,i}. \text{pk}, e(X_{*,1}, Y_*), \Lambda_{*,i}.t_1), \Lambda_{*,i}.t_2)$, then it runs $\text{Ext}(\lambda)$ on \mathcal{A} in order to extract the witness w for the proof $\Lambda_{*,i}.t_2$ and sets $\mathcal{L}[\Lambda_{*,i}.t] \leftarrow w$. Else, G_3 aborts and returns 0. If $g_1^{\mathcal{L}[\Lambda_{*,i}.t]} \neq \Lambda_{*,i}. \text{pk}$ or $e(X_{*,1}, Y_*)^{\mathcal{L}[\Lambda_{*,i}.t]} \neq \Lambda_{*,i}.t_1$, then G_3 aborts and returns 0.

The two games only differ if the extractor Ext fails on a valid proof of knowledge generated by \mathcal{A} , in any of the (at most) q_r calls to the extractor. Hence:

$$|\mathbb{P}[\mathcal{A} \text{ wins } G_1] - \mathbb{P}[\mathcal{A} \text{ wins } G_2]| \leq q_r \cdot \epsilon_{\text{NIPoK}}(\lambda).$$

Finally, we show that $\mathbb{P}[\mathcal{A} \text{ wins } G_2] = 0$. Assume that an adversary \mathcal{A} wins G_2 with non-zero probability. We parse $\pi_{O_*}^{j*}.\tau$ as $(m_X, m'_X, (m_Y, \sigma_Y^1), (m'_Y, \sigma_Y^1), \sigma_X, \sigma'_X, \sigma_Y^2)$, m_X as $(X_1 \| X_2 \| \text{ni}_X)$, m_Y as $(Y \| \text{ni}_Y)$, $\pi_{O_*}^{j*}.\text{PID}$ as (A, B) and $\pi_{O_*}^{j*}.\text{AID}$ as $(\Lambda_i)_{i=1}^n$. On the other hand, we parse sst_* as $(\omega_* \| m_{*,X} \| m_{*,Y} \| \sigma_{*,Y}^1 \| \sigma_{*,X} \| \sigma_{*,Y}^2 \| \sigma_{*,O})$, ω_* as $A_* \| B_* \| (\Lambda_{*,i})_{i=1}^n$, $m_{*,X}$ as $X_{*,1} \| X_{*,2} \| \text{ni}_{*,X}$ and $m_{*,Y}$ as $Y_* \| \text{ni}_{*,Y}$. According to the rules of the game G_1 , $\pi_{O_*}^{j*}.\text{sid} = \text{sid}_*$ and $\pi_{O_*}^{j*}.\text{PID} = (A_*, B_*)$ and $\pi_{O_*}^{j*}.\text{AID} = (\Lambda_{*,i})_{i=1}^n$, which implies that we have $X_{*,1} = X_1$, $X_{*,2} = X_2$, $Y_* = Y$, and $\omega_* = A \| B \| (\Lambda_i)_{i=1}^n$. By definition, Extract returns $k = e(\prod_{i=1}^n \Lambda_i.\text{pk}, X_2)^Y$, where $Y = g_2^y$. We recall that $e(X_1, g_2) = e(g_1, X_2)$.

Algorithm Open(pp, sst_* , $(\Lambda_{*,i}.t)_{i=1}^n$, $(\Lambda_{*,i}.\text{PK})_{i=1}^n$) returns:

$$\begin{aligned} k_* &= \prod_{i=1}^n (\Lambda_{*,i}.t_1) = \prod_{i=1}^n e(X_{*,1}, Y_*)^{\mathcal{L}[\Lambda_{*,i}.t]} \\ &= \prod_{i=1}^n e(X_1, Y)^{\mathcal{L}[\Lambda_{*,i}.t]} = \prod_{i=1}^n e(X_1, g_2^y)^{\mathcal{L}[\Lambda_{*,i}.t]} \\ &= \prod_{i=1}^n e(X_1, g_2)^{y \cdot \mathcal{L}[\Lambda_{*,i}.t]} = \prod_{i=1}^n e(g_1, X_2)^{y \cdot \mathcal{L}[\Lambda_{*,i}.t]} \\ &= \prod_{i=1}^n e(g_1^{\mathcal{L}[\Lambda_{*,i}.t]}, X_2)^y = \prod_{i=1}^n e(\Lambda_{*,i}.\text{pk}, X_2)^y \\ &= \prod_{i=1}^n e(\Lambda_i.\text{pk}, X_2)^y = e(\prod_{i=1}^n \Lambda_i.\text{pk}, X_2)^y \\ &= k, \end{aligned}$$

which implies that $k_* = k$ with probability 1, so \mathcal{A} cannot win the game. This concludes the proof.

7 Implementation of LIKE protocol

We provide a proof-of-concept C-implementation of LIKE, which runs the entire primitive from setup to key-recovery. Since the parties are all implemented on the same machine, we omit network exchanges. The implementation features: two authorities, Alice, Bob, and a single operator (since both operators would have to perform the same computations). We measure the time-complexity of various operations during key-exchange and recovery, and provide averaged results in Table 4.

Setup. We use the optimal Ate pairing [35] over a Barreto-Naehrig curve with 128-bit security. Given [24], we use the BN462 curve described in [22], which follows the recommendations of [?] for pairing curve parameters. We use the base points described in [22] as the generators of groups \mathbb{G}_1 and \mathbb{G}_2 .

For the signature scheme, we use ed25519 [12]. The signature of knowledge and the NIZK proof of a discrete logarithm were implemented using Schnorr proofs and the Fiat-Shamir heuristic, while the NIZK proof of the discrete logarithm equality uses the Chaum and Pedersen protocol and the Fiat-Shamir heuristic, see Section 5. We use SHA-256 as our hash function.

We used mcl [1] for our elliptic-curve and pairing computations. For the ed25519 signature scheme and SHA-256 we used openssl. All tests were done on a Ubuntu 20.04.1 (Linux kernel version : 5.4.0-58) machine with an Intel i5-10210U CPU and 8GB RAM.

Results.: Table 4 shows our results, averaged over 5000 protocol runs. The protocol steps involving the authorities (trapdoor generation and opening) have a complexity that is linear in the number of authorities (2 in our case). In the table, A, B, and O denote

the total runtime of Alice, Bob, and the operator respectively, during authenticated key-exchange (omitting the verifications that are considered as pre-computations). Verification, TDGen, and Open represent the total runtime for those protocol steps. However, note that some of those steps could be run in parallel (for instance, all the authorities can generate trapdoors at the same time) yielding only half the values obtained in our table. For the opening algorithm, authorities can verify the NIZK proofs in parallel, but must then share the complexity of recovering the key. For reference we also include the runtime of a single pairing – since it is our most expensive computation.

Operations	A	B	O	Verification	TDGen	Open	Pairing
CPU time (ms)	7.6	12.5	9.2	10.8	15.5	8.3	2.5

Table 4: Average CPU time (in milliseconds) for 5000 trials.

We note that the cryptographic computations Alice and Bob must perform are relatively inexpensive, particularly since the pairing can be done outside the SIM card. Although we have not taken into account the network operations required to transmit messages, these are unlikely to affect the endpoints. The operator, which *would* be affected, is assumed to have the capacity of running a large number of sessions in parallel – as in fact they do even presently.

8 Conclusion

Our paper is motivated by the observation that lawful interception does not imply that mobile users have to forgo their privacy with respect to mobile operators. To reconcile the two requirements, we introduce Lawful-Interception Key Exchange (LIKE), a new primitive that augments user privacy without harming LI. LIKE guarantees that Alice’s and Bob’s secure channel remains secure except with respect to: Alice, Bob, or a collusion of n legitimate authorities. Neither the operator, nor a collusion of less than n maliciously-behaving authorities can break that privacy. In addition, users are guaranteed the impossibility of being framed of wrong-doing, even if all the authorities and operators worked together. Finally, both the operator and the authorities are guaranteed that the LI procedure will reveal the key that Alice and Bob should have computed in any given session.

Our instantiated primitive relies on a digital signature scheme and NIZK proofs and signatures of knowledge for discrete logarithm. We embed the long-term credentials of the authorities on one side of a pairing, Alice and Bob’s ephemeral DH elements on the other, and rely on the hardness of the BDDH problem. Contrary to existing secret-sharing approaches, in our solution the authorities do not need to store shares to Alice’s and Bob’s randomness for each session (which scales badly and is unreasonable given the retroactive effect of LI laws). Instead, the burden of storage rests with the operator, as is currently the case. We require no trusted party; in fact our three properties assume that the users, operators, and authorities may misbehave (in turn or concomitantly).

Our protocol is versatile and its extract-and-open LI mechanism can be used in various configurations of authorities, judges, and operators.

A guarantee not provided by our protocol is against malicious users who run the protocol correctly, but then use a different key (exchanged by other means) to encrypt their communication. This is not a weakness of our scheme, but rather an inevitability as long as parallel means of exchanging keys (not subject to LI) exist.

Another limitation of our work is the fact that it only works for domestic mobile communications, where both operators are subject to the same set of authorities. It is not obvious how to extend our protocol to embed two independent sets of authority keys into the session state without risking two different conversations taking place. This extension is left as future work.

References

1. MCL. <https://github.com/herumi/mcl>, 2020.
2. 3GPP. *TS 33.106 3rd Generation Partnership Project; Technical Specification Group Services and System Aspects; 3G security; Lawful interception requirements (Release 15)*, 6 2018.
3. 3GPP. *TS 33.126 3GPP; Technical Specification Group Services and System Aspects; Security; Lawful Interception requirements (Rel. 16)*, 09 2019.
4. 3GPP. *TS 33.127 3GPP; Technical Specification Group Services and System Aspects; Security; Lawful Interception (LI) architecture and functions (Rel. 16)*, 03 2020.
5. 3GPP. *TS 33.128 3GPP; Technical Specification Group Services and System Aspects; Security; Protocol and procedures for Lawful Interception (LI); Stage 3 (Rel. 16)*, 03 2020.
6. Harold Abelson, Ross Anderson, Steven M. Bellovin, Josh Benaloh, Matt Blaze, Whitfield "Whit" Diffie, John Gilmore, Matthew Green, Susan Landau, Peter G. Neumann, Ronald L. Rivest, Jeffrey I. Schiller, Bruce Schneier, Michael A. Specter, and Daniel J. Weitzner. Keys under doormats. *Communications of the ACM*, 58(10):24–26, 2015.
7. Abdullah Azfar. Implementation and Performance of Threshold Cryptography for Multiple Escrow Agents in VoIP. In *Proceedings of SPIT/IPC*, pages 143–150, 2011.
8. Mihir Bellare and Oded Goldreich. On Defining Proofs of Knowledge. In *CRYPTO '92*, volume 740 of *LNCS*. Springer, 1992.
9. Mihir Bellare and Shafi Goldwasser. Verifiable Partial Key Escrow. In *CCS '97*. ACM, 1997.
10. Mihir Bellare and Ronald L. Rivest. Translucent Cryptography - An Alternative to Key Escrow, and Its Implementation via Fractional Oblivious Transfer. *J. Cryptology*, 12(2), 1999.
11. Mihir Bellare and Phillip Rogaway. Entity authentication and key distribution. In *CRYPTO '93*, volume 773 of *LNCS*. Springer, 1993.
12. Daniel J. Bernstein, Niels Duif, Tanja Lange, Peter Schwabe, and Bo-Yin Yang. High-speed high-security signatures. In *Proceedings of CHES 2011*, pages 124–142, 2011.
13. Xavier Boyen. The uber-assumption family. In *Pairing 2008*. Springer, 2008.
14. Jan Camenisch and Markus Stadler. Efficient group signature schemes for large groups (extended abstract). In *CRYPTO '97*, volume 1294 of *LNCS*. Springer, 1997.
15. Melissa Chase and Anna Lysyanskaya. On signatures of knowledge. In *CRYPTO*, volume 4117 of *LNCS*. Springer, 2006.
16. David Chaum and Torben P. Pedersen. Wallet databases with observers. In *CRYPTO '92*, volume 740 of *LNCS*, pages 89–105. Springer, 1992.
17. Liqun Chen, Dieter Gollmann, and Chris J. Mitchell. Key escrow in mutually mistrusting domains. In *Proceedings of Security Protocols*, pages 139–153, 1996.

18. M. Chen. Escrowable identity-based authenticated key agreement in the standard model. In *Chinese Electronics Journal*, volume 43, pages 1954–1962, 10 2015.
19. Dorothy E. Denning and Dennis K. Branstad. A taxonomy for key escrow encryption systems. *Commun. ACM*, 39(3), 1996.
20. Yvo Desmedt. Abuses in Cryptography and How to Fight Them. In *CRYPTO '88*, volume 403 of *LNCS*. Springer, 1988.
21. Qiang Fan, Mingjian Zhang, and Yue Zhang. Key Escrow Scheme with the Cooperation Mechanism of Multiple Escrow Agents. 2012.
22. IETF. Pairing-friendly curves. <https://datatracker.ietf.org/doc/draft-irtf-cfrg-pairing-friendly-curves/>, 2020.
23. Joe Kilian and Frank Thomson Leighton. Fair Cryptosystems, Revisited: A Rigorous Approach to Key-Escrow (Extended Abstract). In *CRYPTO '95*, volume 963 of *LNCS*. Springer, 1995.
24. Taechan Kim and Razvan Barbulescu. Extended tower number field sieve: A new complexity for the medium prime case. In *Advances in Cryptology – CRYPTO 2016*, volume 9814 of *LNCS*, pages 543–571, 2016.
25. Yu Long, Zhenfu Cao, and Kefei Chen. A dynamic threshold commercial key escrow scheme based on conic. *Appl. Math. Comput.*, 171(2):972–982, 2005.
26. Yu Long, Kefei Chen, and Shengli Liu. Adaptive Chosen Ciphertext Secure Threshold Key Escrow Scheme from Pairing. *Informatica, Lith. Acad. Sci.*, 17(4):519–534, 2006.
27. Keith M. Martin. Increasing Efficiency of International Key Escrow in Mutually Mistrusting Domains. In *Cryptography and Coding*, volume 1355 of *LNCS*, pages 221–232. Springer, 1997.
28. Silvio Micali. Fair Public-Key Cryptosystems. In *CRYPTO '92*, volume 740 of *LNCS*. Springer, 1992.
29. Ilya Mironov and Noah Stephens-Davidowitz. Cryptographic Reverse Firewalls. In *EURO-CRYPT*, volume 9057, pages 657–686. Springer, 2015.
30. Crypto Museum. Clipper chip. Available at <https://www.cryptomuseum.com/crypto/usa/clipper.htm>.
31. Liang Ni, Gongliang Chen, and Jianhua Li. Escrowable identity-based authenticated key agreement protocol with strong security. *Comput. Math. Appl.*, 65(9):1339–1349, 2013.
32. Claus-Peter Schnorr. Efficient identification and signatures for smart cards. In *CRYPTO '89*, volume 435 of *LNCS*, pages 239–252. Springer, 1989.
33. Adi Shamir. Partial key escrow: A new approach to software key escrow. Presented at Key Escrow Conference, 1995.
34. Adi Shamir. Identity-Based Cryptosystems and Signature Schemes. In *CRYPTO*, pages 47–53, 1984.
35. Frederik Vercauteren. Optimal pairings. *IEEE Transactions on Information Theory*, 56(1):455–461, 2010.
36. Zhen Wang, Zhaofeng Ma, Shoushan Luo, and Hongmin Gao. Key escrow protocol based on a tripartite authenticated key agreement and threshold cryptography. *IEEE Access*, 7:149080–149096, 2019.
37. Charles V. Wright and Mayank Varia. Crypto crumple zones: Enabling limited access without mass surveillance. In *Proceedings of EuroS&P 2018*. IEEE, 2018.
38. Adam L. Young and Moti Yung. Kleptography from Standard Assumptions and Applications. In *Proceedings of SCN*, pages 271–290, 2010.